








CCA
Common Component Architecture

A Look at More Complex Component-Based Applications

CCA Forum Tutorial Working Group
<http://www.cca-forum.org/tutorials/>
tutorial-wg@cca-forum.org

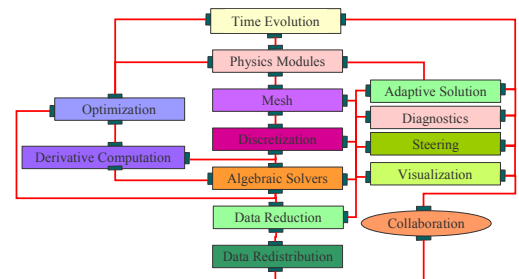


CCA
Common Component Architecture

Complex Applications

Modern Scientific Software Development


- Terascale computing will enable high-fidelity calculations based on multiple coupled physical processes and multiple physical scales
 - Adaptive algorithms and high-order discretization strategies
 - Composite or hybrid solution strategies
 - Sophisticated numerical tools



```

graph TD
    TE[Time Evolution] --> PM[Physics Modules]
    PM --> M[Mesh]
    M --> D[Discretization]
    D --> AS[Algebraic Solvers]
    AS --> DR[Data Reduction]
    DR --> DRD[Data Redistribution]
    DRD --> TE
    AS --> Opt[Optimization]
    AS --> DC[Derivative Computation]
    AS --> ASol[Adaptive Solution]
    AS --> Diag[Diagnostics]
    AS --> Steer[Steering]
    AS --> Viz[Visualization]
    AS --> Coll[Collaboration]
    
```

2




CCA
Common Component Architecture

Complex Applications

Overview

- Using components in high performance simulation codes
 - Examples of increasing complexity
 - Performance
 - Single processor
 - Scalability
- Developing components for high performance simulation codes
 - Strategies for thinking about your own application
 - Developing interoperable and interchangeable components

3



CCA
Common Component Architecture

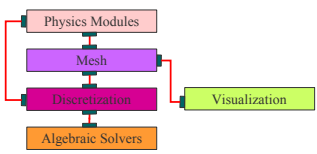
Complex Applications

Our Starting Point

$$\nabla^2 \phi(x, y) = 0 \quad (x, y) \in [0, 1] \times [0, 1]$$

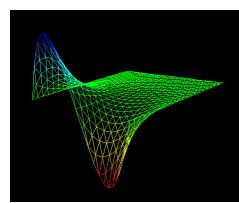
$$\phi(0, y) = 0 \quad \phi(1, y) = \sin(2\pi y)$$

$$\delta\phi/\delta y(x, 0) = \delta\phi/\delta y(x, 1) = 0$$




```

graph LR
    PM[Physics Modules] --> M[Mesh]
    M --> D[Discretization]
    D --> AS[Algebraic Solvers]
    AS --> V[Visualization]
    
```



4




CCA
Common Component Architecture

Complex Applications

Numerical Solution of Example 1

- Physics: Poisson's equation
- Grid: Unstructured triangular mesh
- Discretization: Finite element method
- Algebraic Solvers: PETSc (Portable Extensible Toolkit for Scientific Computation)
- Visualization: VTK tool
- Original Language: C

5




CCA
Common Component Architecture

Complex Applications

Creating Components: Step 1

- Separate the application code into well-defined pieces that encapsulate functionalities
 - Decouple code along numerical functionality
 - Mesh, discretization, solver, visualization
 - Physics is kept separate
 - Determine what questions each component can ask of and answer for other components (this determines the ports)
 - Mesh provides geometry and topology (needed by discretization and visualization)
 - Mesh allows user defined data to be attached to its entities (needed by physics and discretization)
 - Mesh *does not* provide access to its data structures
 - If this is not part of the original code design, this is by far the hardest, most time-consuming aspect of componentization

6




CCA
Common Component Architecture

Complex Applications

Creating the Components: Step 2

- Writing C++ Components
 - Create an abstract base class for each port
 - Create C++ objects that inherit from the abstract base port class and the CCA component class
 - Wrap the existing code as a C++ object
 - Implement the setServices method
- This process was significantly less time consuming (with an expert present) than the decoupling process
 - Lessons learned
 - Definitely look at an existing, working example for the targeted framework
 - Experts are very handy people to have around ;-)

7

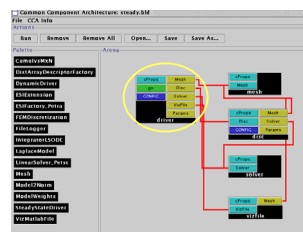


CCA
Common Component Architecture

Complex Applications

The Componentized Example

- The Driver Component
 - Responsible for the overall application flow
 - Initializes the mesh, discretization, solver and visualization components
 - Sets the physics parameters and boundary condition information



8

Complex Applications

9

- [illegible]

0

Complex Applications

[illegible]

11

- [illegible]

2

CCA
Common Component Architecture

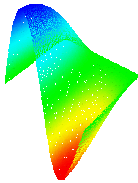
Complex Applications

The next step... time dependence

$$\delta\phi/\delta t = \nabla^2\phi \quad (x,y,t) \in [0,1] \times [0,1]$$

$$\phi(0,y,t)=0 \quad \phi(1,y,t)=.5\sin(2\pi y)\cos(t/2)$$

$$\delta\phi/\delta y(x,0) = \delta\phi/\delta y(x,1) = 0$$

$$\phi(x,y,0)=\sin(.5\pi x) \sin(2\pi y)$$


```

graph TD
    TE[Time Evolution] --> PM[Physics Modules]
    PM --> M[Mesh]
    M --> D[Discretization]
    D --> AS[Algebraic Solvers]
    AS --> DA[Distributed Arrays]
    DA --> DR[Data Redistribution]
    DR --> V[Visualization]
    V --> TE
    
```

13

CCA
Common Component Architecture

Complex Applications

Some things change...

- Requires a time integration component
 - Based on the LSODE library (LLNL)
 - Component implementation developed by Ben Allan (SNL)
- Uses a new visualization component
 - Based on AVS
 - Requires an MxN data redistribution component
 - Developed by Jim Kohl (ORNL)
- The MxN redistribution component requires a Distributed Array Descriptor component
 - Similar to HPF arrays
 - Developed by David Bernholdt (ORNL)
- The driver component changes to accommodate the new physics

14

CCA
Common Component Architecture

Complex Applications

... and some things stay the same

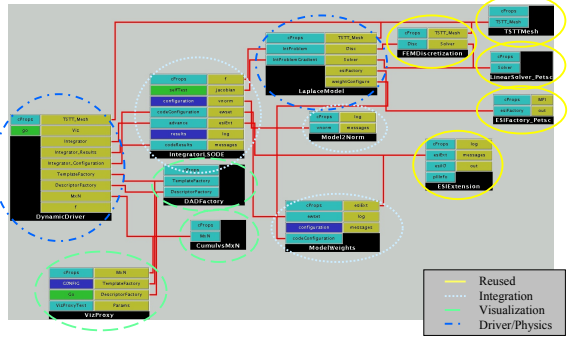
- The mesh component doesn't change
- The discretization component doesn't change
- The solver component doesn't change
 - What we use from the solver component changes
 - Only vectors are needed

15


CCA
Common Component Architecture

Complex Applications

The CCA wiring diagram



16



CCA


Common Component Architecture

Complex Applications

What did this exercise teach us?

- It was easy to incorporate the functionalities of components developed at other labs and institutions given a well-defined interface and header file.
 - In fact, some components (one uses and one provides) were developed simultaneously across the country from each other after the definition of a header file.
 - Amazingly enough, they usually “just worked” when linked together (and debugged individually).
- In this case, the complexity of the component-based approach was higher than the original code complexity.
 - Partially due to the simplicity of this example
 - Partially due to the limitations of some of the current implementations of components

17



CCA

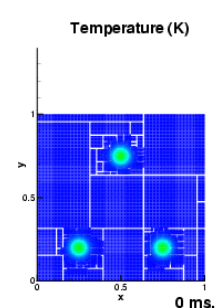
Common Component Architecture

Complex Applications

Beyond the heat equation...

- Flame Approximation
 - H₂-Air mixture; ignition via 3 hot-spots
 - 9-species, 19 reactions, stiff chemistry
- Governing equation


$$\frac{\partial Y_i}{\partial t} = \nabla \cdot \alpha \nabla Y_i + \dot{w}_i$$
- Domain
 - 1cm X 1cm domain
 - 100x100 coarse mesh
 - finest mesh = 12.5 micron.
- Timescales
 - O(10ns) to O(10 microseconds)



Temperature (K)

0 ms.

18



CCA


Common Component Architecture

Complex Applications

Numerical Solution

- Adaptive Mesh Refinement: GrACE
- Stiff integrator: CVODE (LLNL)
- Diffusive integrator: 2nd Order Runge Kutta
- Chemical Rates: legacy f77 code (SNL)
- Diffusion Coefficients: legacy f77 code (SNL)
- New code less than 10%

19




CCA

Common Component Architecture

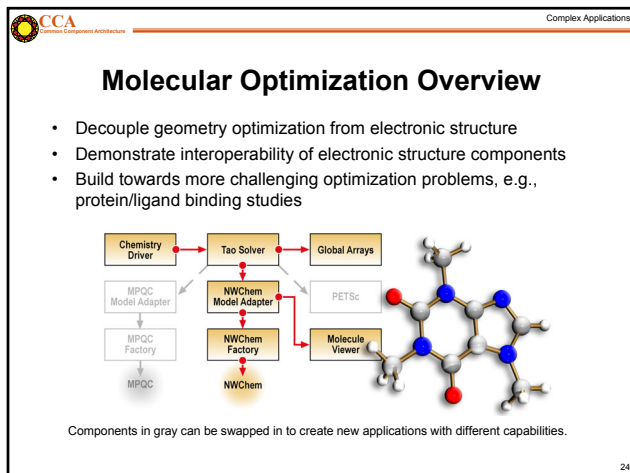
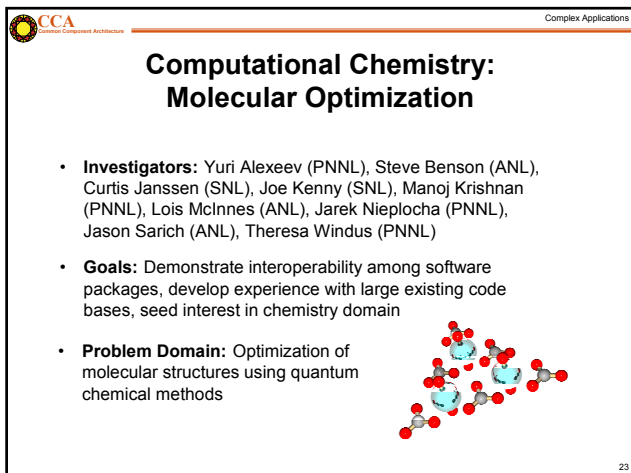
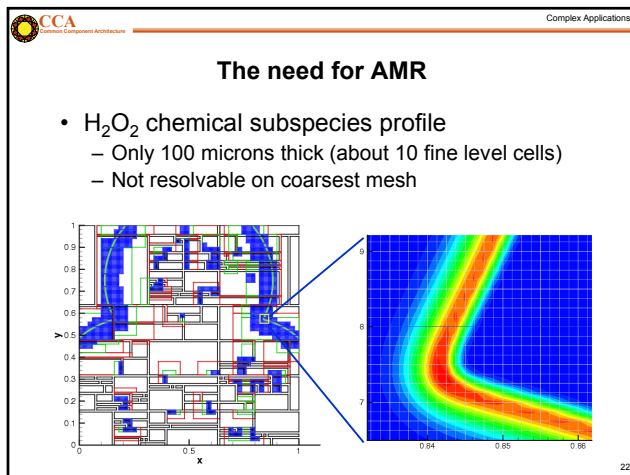
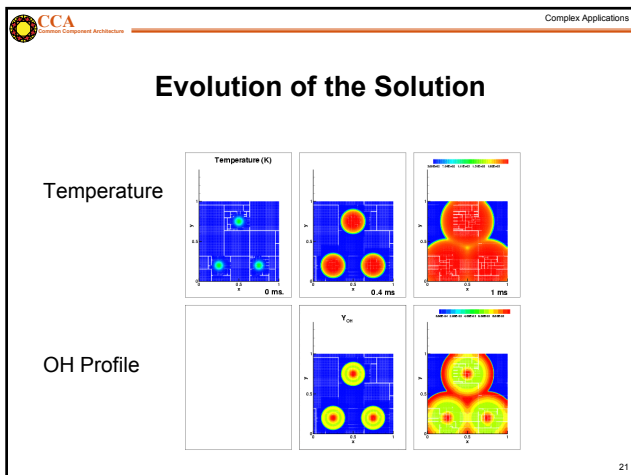
Complex Applications

The CCA Wiring Diagram



— Reused
— Slow Time Scale Integration
— Fast Time Scale Integration
— Driver/Physics

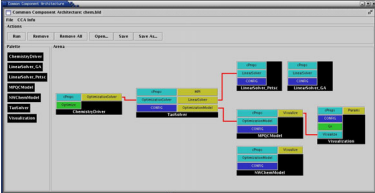
20



CCA
Common Component Architecture

Complex Applications

Wiring Diagram for Molecular Optimization



- Electronic structures components:
 - MPQC (SNL) <http://eros.ca.sandia.gov/~cljanss/mpqc>
 - NWChem (PNNL) <http://www.emsl.pnl.gov/pub/docs/nwchem>
- Optimization components: TAO (ANL) <http://www.mcs.anl.gov/tao>
- Linear algebra components:
 - Global Arrays (PNNL) <http://www.emsl.pnl.gov:2080/docs/global/ga.html>
 - PETSc (ANL) <http://www.mcs.anl.gov/petsc>

25

CCA
Common Component Architecture

Complex Applications

Molecular Optimization Summary

- CCA Impact
 - Demonstrated unprecedented interoperability in a domain not known for it
 - Demonstrated value of collaboration through components
 - Gained experience with several very different styles of "legacy" code
- Future Plans
 - Extend to more complex optimization problems
 - Extend to deeper levels of interoperability

26

CCA
Common Component Architecture

Complex Applications

Componentized Climate Simulations

- NASA's ESMF project has a component-based design for Earth system simulations
 - ESMF components can be assembled and run in CCA compliant frameworks such as Ccaffeine.
- Zhou et al (NASA Goddard) has integrated a simple coupled Atmosphere-Ocean model into Ccaffeine and is working on the Cane-Zebiak model, well-known for predicting *El Nino* events.
- Different PDEs for ocean and atmosphere, different grids and time-stepped at different rates.
 - Synchronization at ocean-atmosphere interface; essentially, interpolations between meshes
 - Ocean & atmosphere advanced in sequence
- Intuitively : Ocean, Atmosphere and 2 coupler components
 - 2 couplers : atm-ocean coupler and ocean-atm coupler.
 - Also a Driver / orchestrator component.

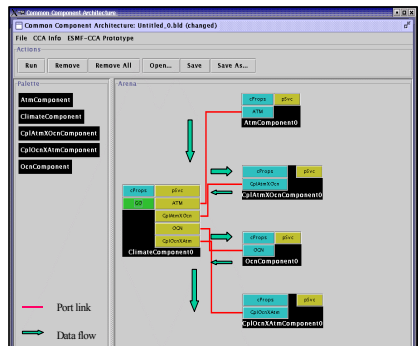
27

CCA
Common Component Architecture

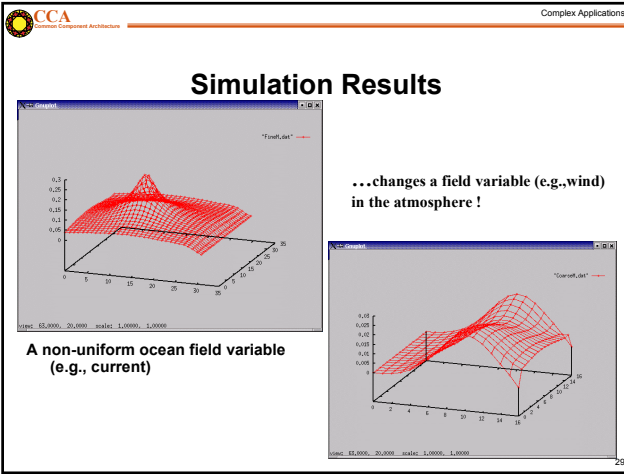
Complex Applications

Coupled Atmosphere-Ocean Model Assembly

- Climate Component :**
 - Schedule component coupling
- Data flow is via pointer NOT data copy.**
 - All components in C++; run in CCAFFEINE.
- Multiple ocean models with the same interface**
 - Can be selected by a user at runtime



28



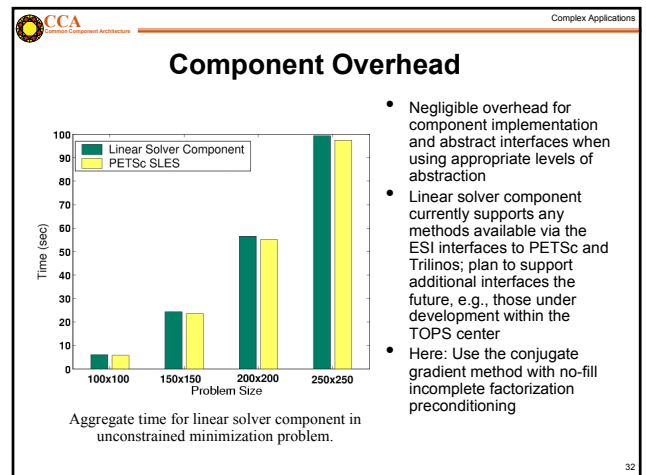
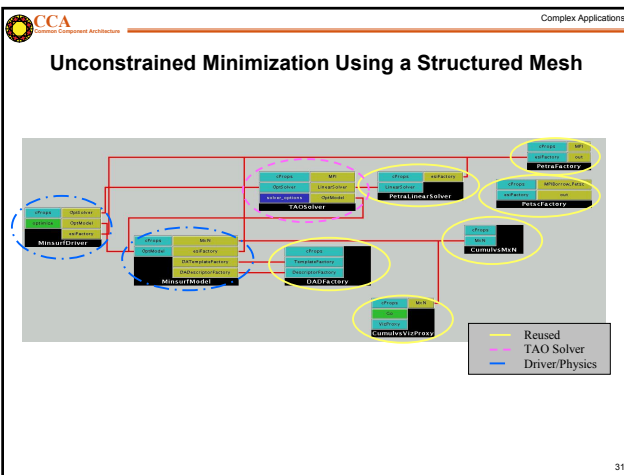
CCA
Common Component Architecture

Complex Applications

Unconstrained Minimization Problem

- Given a rectangular 2-dimensional domain and boundary values along the edges of the domain
- Find the surface with minimal area that satisfies the boundary conditions, i.e., compute $\min f(x)$, where $f: \mathbb{R} \rightarrow \mathbb{R}$
- Solve using optimization components based on TAO (ANL)

30



CCA

Common Component Architecture

Complex Applications

Overhead from Component Invocation

- Invoke a component with different arguments
 - Array
 - Complex
 - Double Complex
- Compare with f77 method invocation
- Environment
 - 500 MHz Pentium III
 - Linux 2.4.18
 - GCC 2.95.4-15
- Components took 3X longer
- Ensure granularity is appropriate!
- Paper by Bernholdt, Elwasif, Kohl and Epperly

Function arg type	f77	Component
Array	80 ns	224ns
Complex	75ns	209ns
Double complex	86ns	241ns

33

CCA

Common Component Architecture

Complex Applications

Scalability on a Linux Cluster

Number of Processors	Components Time (sec)	TAO Time (sec)
1	~110	~110
2	~60	~60
4	~30	~30
8	~15	~15
16	~8	~8

- Newton method with line search
- Solve linear systems with the conjugate gradient method and block Jacobi preconditioning (with no-fill incomplete factorization as each block's solver, and 1 block per process)
- Negligible component overhead; good scalability

Total execution time for the minimum surface minimization problem using a fixed-sized 250x250 mesh.

34

CCA

Common Component Architecture

Complex Applications

List of Component Re-Use

- Various services in Ccaffeine
- Integrator
 - IntegratorLSODE (2)
 - RK2 (2)
- Linear solvers
 - LinearSolver_Petra (4)
 - LinearSolver_PETSc (4)
- AMR
 - AMRmesh (3)
- Data description
 - DADFactory (3)
- Data redistribution
 - CumulvsMxN (3)
- Visualization
 - CumulvsVizProxy (3)

Component interfaces to numerical libraries

Component interfaces to parallel data management and visualization tools

35

CCA

Common Component Architecture

Complex Applications

The Next Level

- Common Interface Specification
 - Provides plug-and-play interchangeability
 - Requires domain specific experts
 - Typically a difficult, time-consuming task
 - A success story: MPI
- A case study... the TSTT/CCA mesh interface
 - TSTT = Terascale Simulation Tools and Technologies (www.tstt-scidac.org)
 - A DOE SciDAC ISIC focusing on meshes and discretization
 - Goal is to enable
 - hybrid solution strategies
 - high order discretization
 - Adaptive techniques

36

CCA
Common Component Architecture

Complex Applications

Current Situation

Current Situation

- Public interfaces for numerical libraries are unique
- Many-to-Many* couplings require *Many*² interfaces
 - Often a heroic effort to understand the inner workings of both codes
- Not a scalable solution

37

CCA
Common Component Architecture

Complex Applications

Common Interface Specification

Reduces the *Many-to-Many* problem to a *Many-to-One* problem

- Allows interchangeability and experimentation
- Challenges
 - Interface agreement
 - Functionality limitations
 - Maintaining performance

38

CCA
Common Component Architecture

Complex Applications

TSTT Philosophy

Create a small set of interfaces that existing packages can support

AOMD, CUBIT, Overture, GrACE, ...

Enable both interchangeability and interoperability

Balance performance and flexibility

Work with a large tool provider and application community to ensure applicability

Tool providers: TSTT and CCA SciDAC centers

Application community: SciDAC and other DOE applications

39


CCA
Common Component Architecture

Complex Applications

Basic Interface

- Enumerated types
 - Entity Type: VERTEX, EDGE, FACE, REGION
 - Entity Topology: POINT, LINE, POLYGON, TRIANGLE, QUADRILATERAL, POLYHEDRON, TETRAHEDRON, HEXAHEDRON, PRISM, PYRAMID, SEPTAHEDRON
- Opaque Types
 - Mesh, Entity, Workset, Tag
- Required interfaces
 - Entity queries (geometry, adjacencies), Entity iterators, Array-based query, Workset iterators, Mesh/Entity Tags, Mesh Services

40



CCA


Common Component Architecture

Complex Applications

Issues that have arisen

- Nomenclature is harder than we first thought
- Cannot achieve the 100 percent solution, so...
 - What level of functionality should be supported?
 - Minimal interfaces only?
 - Interfaces for convenience and performance?
 - What about support of existing packages?
 - Are there atomic operations that all support?
 - What additional functionalities from existing packages should be required?
 - What about additional functionalities such as locking?
- Language interoperability is a problem
 - Most TSTT tools are in C++, most target applications are in Fortran
 - How can we avoid the "least common denominator" solution?
 - Exploring the SIDL/Babel language interoperability tool

41



CCA


Common Component Architecture

Complex Applications

Summary

- Complex applications that use components are possible
 - Combustion
 - Chemistry applications
 - Optimization problems
 - Climate simulations
- Component reuse is significant
 - Adaptive Meshes
 - Linear Solvers (PETSc, Trilinos)
 - Distributed Arrays and MxN Redistribution
 - Time Integrators
 - Visualization
- Examples shown here leverage and extend parallel software and interfaces developed at different institutions
 - Including CUMULVS, ESI, GrACE, LSODE, MPICH, PAWS, PETSc, PVM, TAO, Trilinos, TSTT.
- Performance is not significantly affected by component use
- Definition of domain-specific common interfaces is key

42



CCA

Common Component Architecture

Complex Applications

Componentizing your own application

- The key step: think about the decomposition strategy
 - By physics module?
 - Along numerical solver functionality?
 - Are there tools that already exist for certain pieces? (solvers, integrators, meshes?)
 - Are there common interfaces that already exist for certain pieces?
 - Be mindful of the level of granularity
- Decouple the application into pieces
 - Can be a painful, time-consuming process
- Incorporate CCA-compliance
- Compose your new component application
- Enjoy!

43