



Common Component Architecture Concepts

CCA Forum Tutorial Working Group

<http://www.cca-forum.org/tutorials/>
tutorial-wg@cca-forum.org



Goals

- Introduce the **motivation** and essential **features** of the Common Component Architecture
- Provide **common vocabulary** for remainder of tutorial
- What distinguishes CCA from **other component environments**?

Historical Perspective

Scientific Computing

- Has focused on extracting **performance**, **parallelism**, etc.
- Code **complexity has increased** due to both performance and scientific issues
- Little has been done to address the complexity

Component Models

- Have focused on **managing complexity**
- Commodity component models are used primarily in the **business** and **internet** software arenas
- Performance (at HPC level) not addressed

CCA is intended to bring the benefits of components to scientific computing

3

What is the CCA? (User View)

- A component model specifically designed for **high-performance** scientific computing
- Supports both **parallel and distributed** applications
- Designed to be implementable **without sacrificing performance**
- **Minimalist** approach makes it easier to componentize existing software


4

What is the CCA? (2)

- Components are **peers**
- *Not* just a dataflow model
- A **tool** to enhance the productivity of scientific programmers
 - Make the hard things easier, make some intractable things tractable
 - Support & promote reuse & interoperability
 - **Not a magic bullet**

5


CCA Concepts: Ports



- Components interact through well-defined **interfaces**, or **ports**
 - In OO languages, a port is a class or interface
 - In Fortran, a port is a bunch of subroutines or a module
- Components may **provide** ports – **implement** the class or subroutines of the port (**"Provides" Port**)
- Components may **use** ports – **call** methods or subroutines in the port (**"Uses" Port**)
- Links denote a procedural (caller/callee) relationship, **not dataflow!**
 - e.g., FunctionPort could contain: *evaluate*(*in* Arg, *out* Result)

6

CCA Concepts: Components



- Components are a unit of software **composition**
- Components provide/use one or more **ports**
 - A component with no ports isn't very interesting
- Components include some code which **interacts with the CCA framework**
 - Implement `setServices` method, constructor, destructor
 - Use `getPort/releasePort` to access ports on other components
- The **granularity** of components is dictated by the application architecture and by performance considerations
- Components are **peers**
 - Application architecture determines relationships

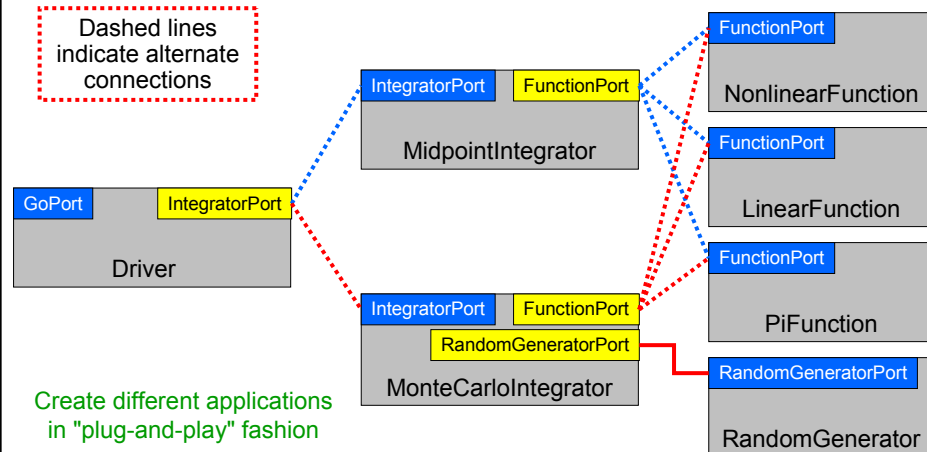
7

CCA Concepts: Frameworks

- The framework provides the means to “hold” components and **compose** them into applications
 - The framework is often application's “main” or “program”
- Frameworks allow **exchange of ports** among components without exposing implementation details
- Frameworks provide a small set of **standard services** to components
 - BuilderService allow programs to compose CCA apps
- Frameworks may make themselves **appear as components** in order to connect to components in other frameworks
- *Currently:* specific frameworks support specific computing models (parallel, distributed, etc.).
Future: full flexibility through integration or interoperation

8

Components and Ports in the Integrator Example



9

Writing Components

- Components...
 - Inherit from **gov.cca.Component**
 - Implement **setServices** method to register ports this component will **provide** and **use**
 - Implement the ports they they provide
 - Use ports on other components
 - **getPort/releasePort** from framework **Services** object
- Interfaces (ports) extend **gov.cca.Port**

Much more detail later in the tutorial!

10

Adapting Existing Code into Components

- Suitably structured code (programs, libraries) should be relatively easy to adapt to CCA
- Decide **level of componentization**
 - Can evolve with time (start with coarse components, later refine into smaller ones)
- Define **interfaces** and write wrappers between them and existing code
- Add **framework interaction code** for each component
 - **setServices**, constructor, destructor
- Modify component internals to **use other components** as appropriate
 - **getPort**, **releasePort** and method invocations

11

Writing Frameworks

- ***There is no reason for most people to write frameworks – just use the existing ones!***
- Frameworks must provide certain ports...
 - **ConnectionEventService**
 - Informs the component of connections
 - **AbstractFramework**
 - Allows the component to *behave as a framework*
 - **BuilderService**
 - instantiate components & connect ports
 - **ComponentRepository**
 - A default place where components are found
 - Coming soon: framework services can be implemented in components and registered as services
- Frameworks must be able to load components
 - Typically shared object libraries, can be statically linked
- Frameworks must provide a way to compose applications from components

12

Typical Component Lifecycle

- **Composition Phase**

- Component is **instantiated** in framework
- Component interfaces are **connected** appropriately

We'll look at actual code in next tutorial module

- **Execution Phase**

- Code in components uses functions provided by another component

- **Decomposition Phase**

- **Connections** between component interfaces may be **broken**
- Component may be **destroyed**


In an application, individual components may be in different phases at different times

Steps may be under human or software control

13

User Viewpoint: Loading and Instantiating Components

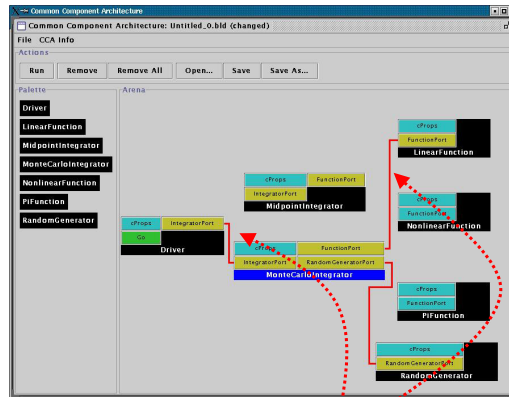
- Components are code (usu. library or shared object) + metadata
- Using metadata, a **Palette** of available components is constructed
- Components are instantiated by user action (i.e. by dragging from **Palette** into **Arena**)
- Framework calls component's **constructor**, then **setServices**
- Details are **framework-specific!**
- **Ccaffeine** currently provides both command line and GUI approaches



14

User Connects Ports

- Can only connect uses & provides
 - Not uses/uses or provides/provides
- Ports connected by type, not name
 - Port names must be unique within component
 - Types must match across components
- Framework puts info about *provider* of port into *using component's* Services object



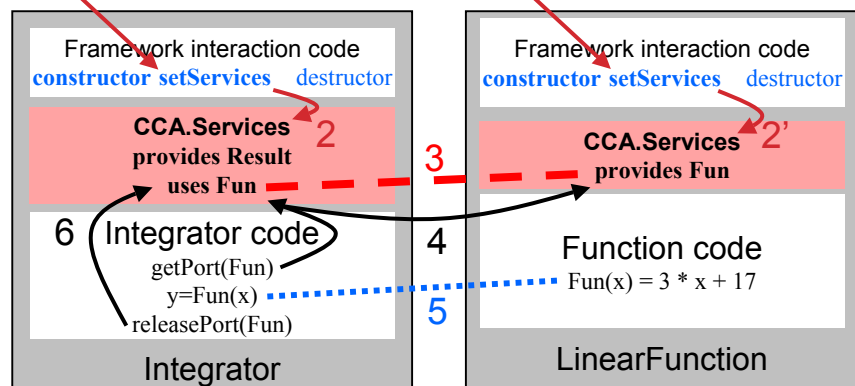
connect	Driver	IntegratorPort	MonteCarloIntegrator	IntegratorPort
connect	MonteCarloIntegrator	FunctionPort	LinearFunction	FunctionPort
...				

15

Framework Mediates *Most** Component Interactions

* Method invocation need not be mediated by the framework!


Execution Phase




16

Component's View of Instantiation

- Framework calls component's **constructor**
- Component initializes internal data, etc.
 - Knows *nothing* outside itself

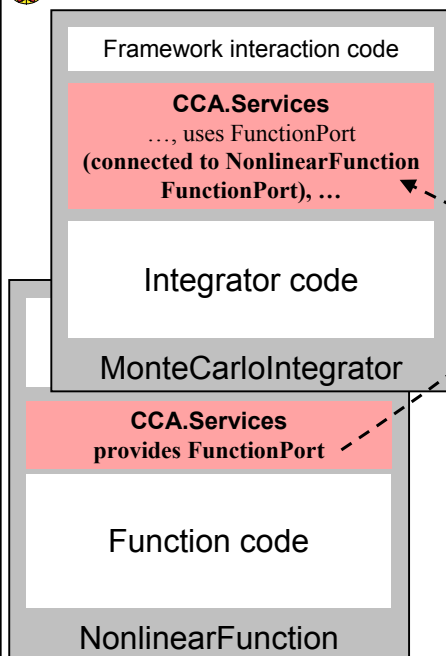


- Framework calls component's **setServices**
 - Passes setServices an object representing everything "outside"
 - setServices declares ports component *uses* and *provides*
- Component *still* knows nothing outside itself
 - But Services object provides the means of communication w/ framework
- Framework now knows how to "decorate" component and how it might connect with others



17

Component's View of Connection

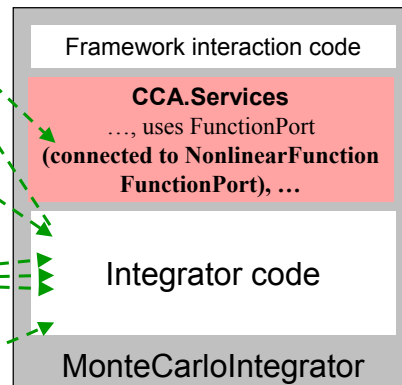


- Framework puts info about provider into **user component's Services** object
 - **MonteCarloIntegrator's** Services object is aware of connection
 - **NonlinearFunction** is not!
- **MCI's** integrator code cannot yet call functions on FunctionPort

18

Component's View of Using a Port


- User calls `getPort` to obtain (handle for) port from Services
 - Finally user code can “see” provider
- `Cast` port to expected type
 - OO programming concept
 - Insures type safety
 - Helps enforce declared interface
- `Call` methods on port
 - e.g.
 $\text{sum} = \text{sum} + \text{function} \rightarrow \text{evaluate}(x)$
- `Release` port



19

Importance of Provides/Uses Pattern for Ports

- Fences between components
 - Components must `declare` both what they provide and what they use
 - Components `cannot interact` until ports are connected
 - No mechanism to call anything not part of a port
- Ports preserve high performance `direct connection` semantics...
- ...While also allowing `distributed computing`



20

CCA Concepts: “Direct Connection” Maintains Local Performance

- Calls *between* components equivalent to a C++ *virtual function call*: lookup function location, invoke it
 - Cost equivalent of *~2.8 F77 or C function calls*
 - *~48 ns vs 17 ns* on 500 MHz Pentium III Linux box
- *Language interoperability* can impose additional overheads
 - Some arguments require conversion
 - Costs vary, but small for typical scientific computing needs
- Calls *within* components have *no CCA-imposed overhead*
- **Implications**
 - *Be aware of costs*
 - Design so inter-component calls *do enough work* that overhead is negligible

21

How Does Direct Connection Work?


- Components loaded into *separate namespaces* in the *same address* space (process) from shared libraries
- *getPort* call returns a pointer to the port's function table
- All this happens “automatically” – *user just sees high performance*
 - *Description reflects Ccaffeine implementation, but similar or identical mechanisms are in other direct connect fwks*
- *Many CORBA implementations offer a similar approach to improve performance, but using it violates the CORBA standards!*

22

CCA Concepts: Framework Stays “Out of the Way” of Component Parallelism

- Single component multiple data (SCMD) model is component analog of widely used SPMD model
- Each process loaded with the same set of components wired the same way
- Different components in same process “talk to each other” via ports and the framework

Same component in different processes talk to each other through their favorite communications layer (i.e. MPI, PVM, GA)



Components: Blue, Green, Red

Framework: Gray


MCMD/MPMD also supported

Other component models ignore parallelism entirely

23

Scalability of Scientific Data Components in CFRFS Combustion Applications

- Investigators: S. Lefantzi, J. Ray, and H. Najm (SNL)
- Uses GrACEComponent, CvodesComponent, etc.
- Shock-hydro code with no refinement
- 200 x 200 & 350 x 350 meshes
- Cplant cluster
 - 400 MHz EV5 Alphas
 - 1 Gb/s Myrinet
- Negligible component overhead
- Worst perf : 73% scaling efficiency for 200x200 mesh on 48 procs




Reference: S. Lefantzi, J. Ray, and H. Najm, Using the Common Component Architecture to Design High Performance Scientific Simulation Codes, *Proc of Int. Parallel and Distributed Processing Symposium*, Nice, France, 2003, accepted.

24

CCA Concepts: MxN Parallel Data Redistribution

- Share Data Among Coupled Parallel Models
 - Disparate Parallel Topologies (M processes vs. N)
 - e.g. Ocean & Atmosphere, Solver & Optimizer...
 - e.g. Visualization (Mx1, increasingly, MxN)





Research area -- tools under development

25

CCA Concepts: Language Interoperability

- Existing language interoperability approaches are “point-to-point” solutions
- Babel provides a unified approach in which all languages are considered **peers**
- Babel used primarily at interfaces



Babel tutorial coming up!

Few other component models support all languages and data types important for scientific computing

26

What the CCA isn't...

- CCA doesn't specify who owns "main"
 - CCA components are peers
 - Up to application to define component relationships
 - "Driver component" is a common design pattern
- CCA doesn't specify a parallel programming environment
 - Choose your favorite
 - Mix multiple tools in a single application
- CCA doesn't specify I/O
 - But it gives you the infrastructure to create I/O components
 - Use of `stdio` may be problematic in mixed language env.
- CCA doesn't specify interfaces
 - But it gives you the infrastructure to define and enforce them
 - CCA Forum supports & promotes "standard" interface efforts
- CCA doesn't require (but does support) separation of algorithms/physics from data

27

What the CCA is...

- CCA is a *specification* for a component environment
 - Fundamentally, a design pattern
 - Multiple "reference" implementations exist
 - Being used by applications
- CCA increases productivity
 - Supports and promotes software interoperability and reuse
 - Provides "plug-and-play" paradigm for scientific software
- CCA offers the flexibility to architect your application as you think best
 - Doesn't dictate component relationships, programming models, etc.
 - Minimal performance overhead
 - Minimal cost for incorporation of existing software
- CCA provides an environment in which domain-specific application frameworks can be built
 - While retaining opportunities for software reuse at multiple levels

28

Concept Review

- **Ports**
 - Interfaces between components
 - Uses/provides model
- **Framework**
 - Allows assembly of components into applications
- **Direct Connection**
 - Maintain performance of local inter-component calls
- **Parallelism**
 - Framework stays out of the way of parallel components
- **MxN Parallel Data Redistribution**
 - Model coupling, visualization, etc.
- **Language Interoperability**
 - Babel, Scientific Interface Definition Language (SIDL)