

 **CCA**
Common Component Architecture


Welcome to the Common Component Architecture Tutorial

SciDAC Conference 2007
29 June 2007

CCA Forum Tutorial Working Group
[http://www.cca-forum.org/tutorials/
tutorial-wg@cca-forum.org](http://www.cca-forum.org/tutorials/tutorial-wg@cca-forum.org)


 This work is licensed under a [Creative Commons Attribution 2.5 License](http://creativecommons.org/licenses/by/2.5/)

1


 **CCA**
Common Component Architecture

Licensing Information

- This tutorial is distributed under the [Creative Commons Attribution 2.5 License](http://creativecommons.org/licenses/by/2.5/)
 - <http://creativecommons.org/licenses/by/2.5/>
- In summary, **you are free:**
 - to copy, distribute, display, and perform the work
 - to make derivative works
 - to make commercial use of the work
- **Under the following conditions:**
 - **Attribution.** You must attribute the work in the manner specified by the author or licensor.
- For any reuse or distribution, you must make clear to others the license terms of this work.
- Any of these conditions can be waived if you get permission from the copyright holder.
- **Your fair use and other rights are in no way affected by the above.**
- **Requested reference:**
 - CCA Forum Tutorial Working Group, Common Component Architecture Tutorial, 2007, <http://www.cca-forum.org/tutorials/>




2

 CCA Common Component Architecture			
<h2>Agenda & Table of Contents</h2>			
Time	Title	Slide No.	Presenter
1:30-1:35pm	Welcome	1	David Bernholdt, ORNL
1:35-2:30pm	What Can Component Technology do for Scientific Computing	6	David Bernholdt, ORNL
	An Intro to Components & the CCA	17	David Bernholdt, ORNL
2:30-3:15pm	CCA Tools & Demonstration	72	Rob Armstrong, SNL
3:15-3:30pm	<i>Break (Room 4-349)</i>		
3:30pm-4:15pm	Language Interoperable CCA Components with Babel	116	Gary Kumfert, LLNL
4:15-5:10pm	Using CCA: Approaches & Experience	143	Jim Kohl, ORNL
5:10-5:15pm	Closing	192	Jim Kohl, ORNL
5:15-5:30pm	Questions and Discussion	n/a	All

No time for a hands-on session today, but the materials and codes are available for self-study at <http://www.cca-forum.org/tutorials/>

3

 CCA Common Component Architecture			
<h2>Network Access</h2>			
<ul style="list-style-type: none">• http://web.mit.edu/ist/topics/network/netguests.html• Authorization may take up to 10 minutes• We have no experience with this network, but if you have problems, we'll try to get you some help			

4



Who We Are: The Common Component Architecture (CCA) Forum

- Combination of standards body and user group for the CCA
- Define [specifications](#) for high-performance scientific components & frameworks
- Promote and facilitate development of [tools](#) for component-based software development, [components](#), and component [applications](#)
- Open membership, quarterly meetings...

General mailing list: cca-forum@cca-forum.org

Web: <http://www.cca-forum.org/>

- Center for Technology for Advanced Scientific Component Software (TASCS)
 - Funded by the US DOE SciDAC program
 - Core development team for CCA technologies



5



What Can Component Technology do for Scientific Computing?

CCA Forum Tutorial Working Group

[http://www.cca-forum.org/tutorials/
tutorial-wg@cca-forum.org](http://www.cca-forum.org/tutorials/tutorial-wg@cca-forum.org)



This work is licensed under a [Creative Commons Attribution 2.5 License](http://creativecommons.org/licenses/by/2.5/)

6



Common Component Architecture

Managing Code Complexity

Some Common Situations:

- Your code is so large and complex it has become fragile and hard to keep running
- You have a simple code, and you want to extend its capabilities
 - rationally
- You want to develop a computational “toolkit”
 - Many modules that can be assembled in different ways to perform different scientific calculations
 - Gives users w/o programming experience access to a flexible tool for simulation
 - Gives users w/o HPC experience access to HPC-ready software

How CCA Can Help:

- Components help you think about software in manageable chunks that interact only in well-defined ways
- Components provide a “plug-and-play” environment that allows easy, flexible application assembly

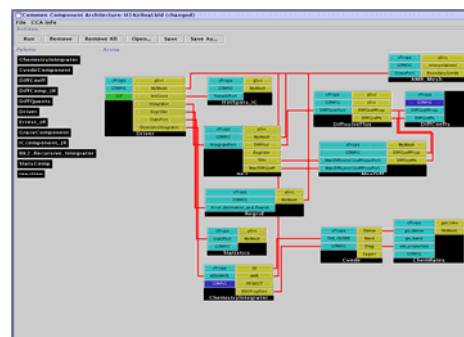
7



Common Component Architecture

Example: Computational Facility for Reacting Flow Science (CFRFS)

- A toolkit to perform simulations of unsteady flames
- Solve the Navier-Stokes with detailed chemistry
 - Various mechanisms up to ~50 species, 300 reactions
 - Structured adaptive mesh refinement
- CFRFS today:
 - 61 components
 - 7 external libraries
 - 9 contributors



“Wiring diagram” for a typical CFRFS simulation, utilizing 12 components.

CCA tools used: Ccaffeine, and ccafe-gui

Languages: C, C++, F77

8

CCA
Common Component Architecture

Helping Groups Work with Software

Some Common Situations:

- Many (geographically distributed) developers creating a large software system
 - Hard to coordinate, different parts of the software don't work together as required
- Groups of developers with different specialties
- Forming communities to standardize interfaces or share code

How CCA Can Help:

- Components are natural units for
 - Expressing software architecture
 - Individuals or small groups to develop
 - Encapsulating particular expertise
- Some component models (including CCA) provide tools to help you think about the *interface* separately from the *implementation*

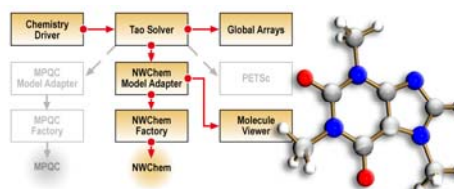


9

CCA
Common Component Architecture

Example: Quantum Chemistry

- Integrated state-of-the-art optimization technology into two quantum chemistry packages to explore effectiveness in chemistry applications
- Geographically distributed expertise:
 - California - chemistry
 - Illinois - optimization
 - Washington – chemistry, parallel data management
- Effective collaboration with minimal face-to-face interaction



Schematic of CCA-based molecular structure determination quantum chemistry application.


Components based on: MPQC, NWChem (quantum chem.), TAO (optimization), Global Arrays, PETSc (parallel linear algebra)

CCA tools used: Babel, Ccaffeine, and ccafe-gui

Languages: C, C++, F77, Python



10



CCA
Common Component Architecture

Example: TSTT Unstructured Mesh Tool Interoperability

- Common interface for unstructured mesh geometry and topology
 - 7 libraries: FMDB, Frontier, GRUMMP, Mesquite, MOAB, NWGrid, Overture
 - 6 institutions: ANL, BNL/SUNY-Stony Brook, LLNL, PNNL, RPI, SNL
- Reduces need for N^2 pairwise interfaces to just N

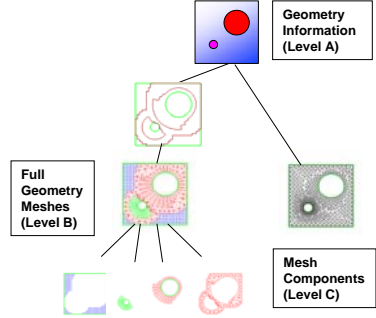



Illustration of geometry domain hierarchy used in TSTT mesh interface.

CCA tools used: Babel (SIDL), Chasm
Library languages: C, C++, F77, F90



CCA
Common Component Architecture


Language Interoperability

Some Common Situations:

- Need to use existing code or libraries written in multiple languages in the same application?
- Want to allow others to access your library from multiple languages?
- Technical or sociological reasons for wanting to use multiple languages in your application?

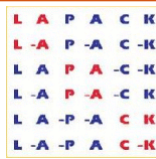
How CCA Can Help:

- Some component models (including CCA) allow transparent mixing of languages
- Babel (CCA's language interop. tool) can be used separately from other component concepts



CCA
Common Component Architecture

Examples



hypre

- High performance preconditioners and linear solvers
- Library written in C
- Babel-generated object-oriented interfaces provided in C, C++, Fortran

"I implemented a Babel-based interface for the hypre library of linear equation solvers. The Babel interface was straightforward to write and gave us interfaces to several languages for less effort than it would take to interface to a single language."


-- Jeff Painter, LLNL. 2 June 2003

LAPACK07

- Update to LAPACK linear algebra library
 - To be released 2007
 - Library written in F77, F95
- Will use Babel-generated interfaces for: C, C++, F77, F95, Java, Python
- Possibly also ScaLAPACK (distributed version)

CCA tools used: Babel, Chasm

13



CCA
Common Component Architecture

Coupling Codes

Some Common Situations:

- Your application makes use of numerous third-party libraries
 - Some of which interact (version dependencies)
- You want to develop a simulation in which your code is coupled with others
 - They weren't designed with this coupling in mind
 - They must remain usable separately too
 - They are all under continual development, individually
 - They're all parallel and need to exchange data frequently

How CCA Can Help:

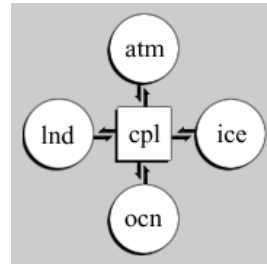
- Components are isolated from one another
 - Interactions via well-defined interfaces
 - An application can include multiple versions of a component
- Components can be composed flexibly, hierarchically
 - Standalone application as one assembly, coupled simulation as another
- CCA can be used in SPMD, MPMD, and distributed styles of parallel computing
- CCA is developing technology to facilitate data and functional coupling of parallel applications

14



Example: Global Climate Modeling and the Model Coupling Toolkit (MCT)

- MCT is the basis for Community Climate System Model (CCSM3.0) coupler (cpl6)
- Computes interfacial fluxes and manages redistribution of data among parallel processes
- Written in F90, Babel-generated bindings for C++, Python
- **CCA tools used:** Babel, Chasm



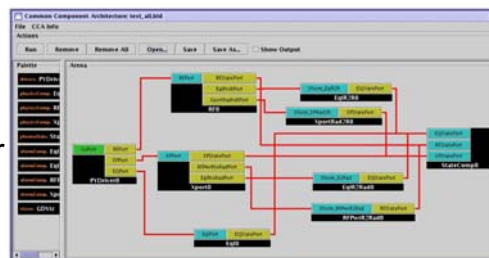
Schematic of CCSM showing coupler managing data exchanges between atmosphere, sea ice, ocean, and land models.
 (From <http://www.cesm.ucar.edu/models/ccsm3.0/cpl6/>)

15



Example: Integrated Fusion Simulation

- Proof-of-principle of using CCA for integrated whole-device modeling needed for the ITER fusion reactor
- Couples radio frequency (RF) heating of plasma with transport modeling
- Coarse-grain encapsulation of pre-existing programs
- Follow-on plans for RF, transport, and magneto-hydrodynamics



"Wiring diagram" for integrated fusion simulation.

Components based on: AORSA, Houlberg's transport library

New components: Driver, State

CCA tools used: Babel, Chasm, Ccaffeine, ccafe-gui

Languages: C++, F90, Python

16



CCA
Common Component Architecture

An Introduction to Components and the Common Component Architecture

CCA Forum Tutorial Working Group

<http://www.cca-forum.org/tutorials/>
tutorial-wg@cca-forum.org



This work is licensed under a [Creative Commons Attribution 2.5 License](https://creativecommons.org/licenses/by/2.5/)

17



CCA
Common Component Architecture

Goals of This Module

- Introduce basic **concepts and vocabulary** of component-based software engineering and the CCA
- Highlight the special **demands of high-performance scientific computing** on component environments
- Give you sufficient **understanding** of the CCA to begin **evaluating** whether it would be useful to you



18



What are Components?

- No universally accepted definition in computer science research, but key features include...
- A unit of software development/deployment/reuse
 - i.e. has **interesting functionality**
 - Ideally, functionality someone else might be able to **(re)use**
 - Can be **developed independently** of other components
- Interacts with the outside world only through well-defined interfaces
 - **Implementation is opaque** to the outside world
- Can be composed with other components
 - “Plug and play” model to build applications
 - **Composition based on interfaces**



19

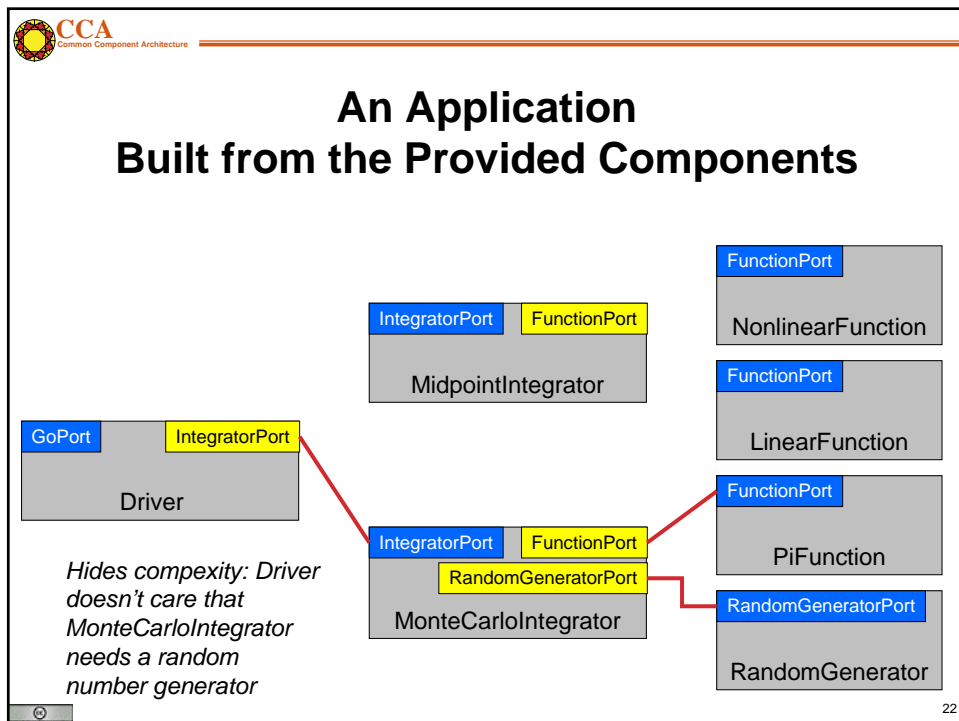
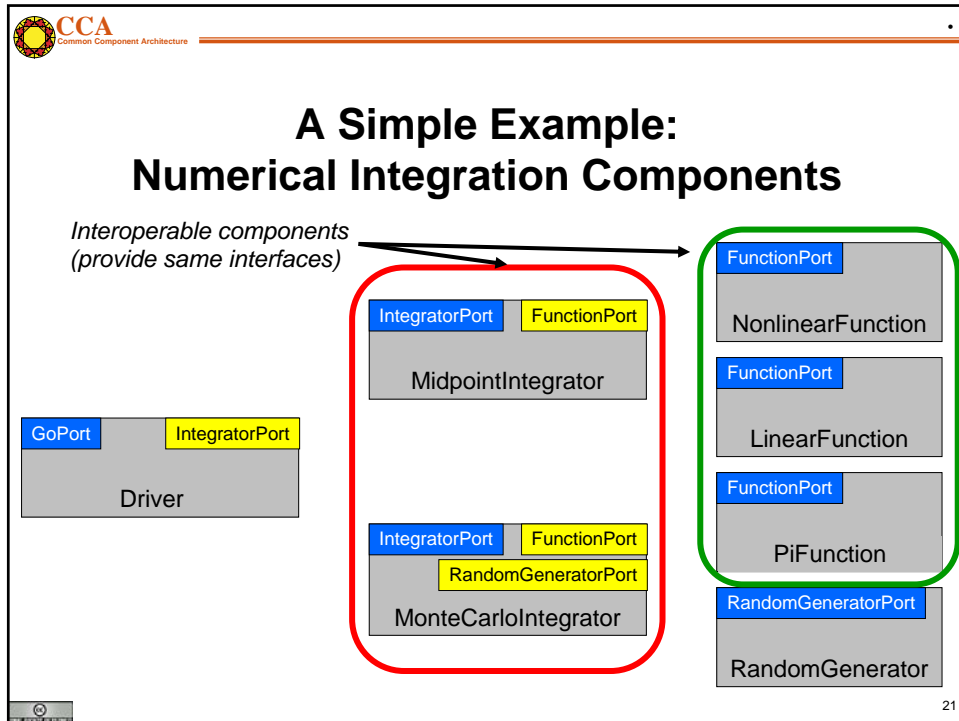


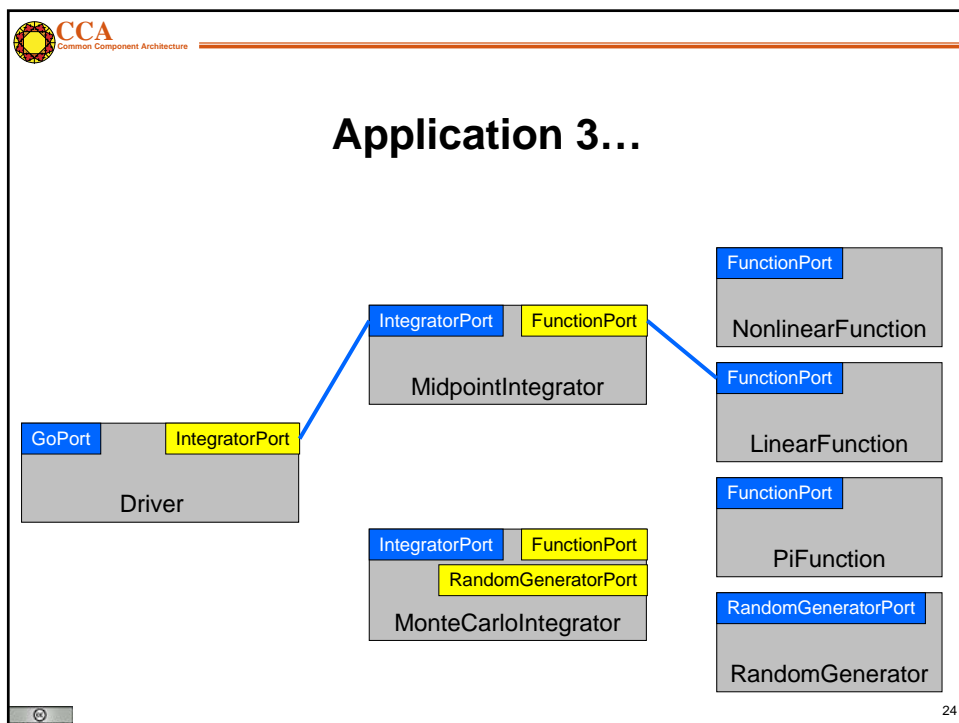
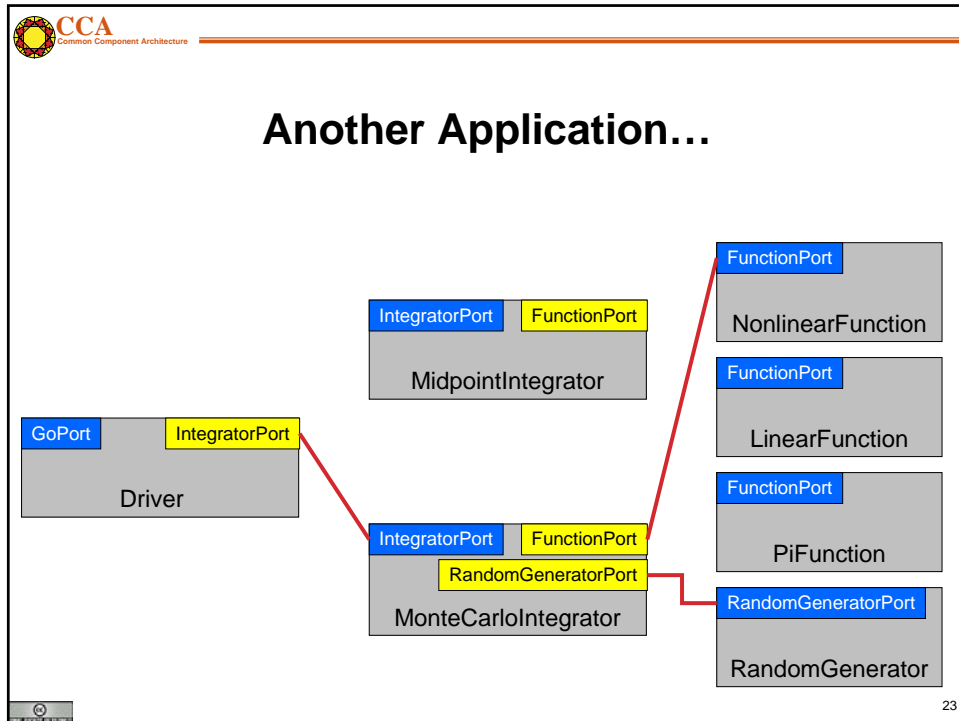
What is a Component Architecture?

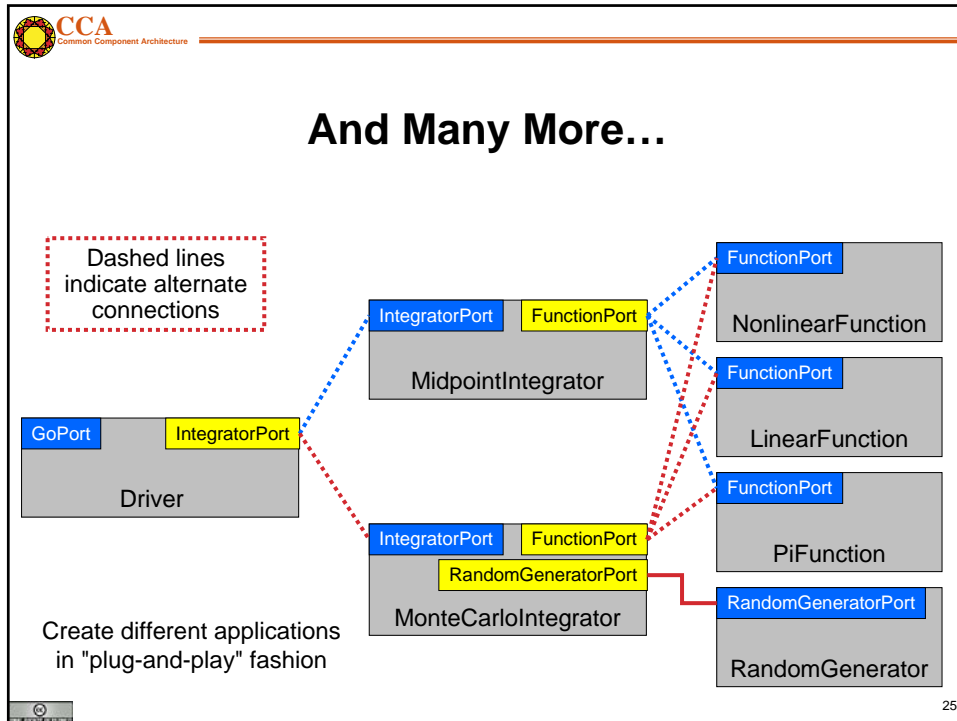
- A set of **standards** that allows:
 - Multiple groups to write units of software (**components**)...
 - And have confidence that their components will **work with other components** written in the same architecture
- These standards **define**...
 - The rights and responsibilities of a **component**
 - How components express their **interfaces**
 - The environment in which components are composed to form an application and executed (**framework**)
 - The rights and responsibilities of the framework



20







CCA
Common Component Architecture

Comparison of Application Development Approaches

Characteristics	Monolithic Simulation Code	Simulation Frameworks	Library -Based	Component -Based
Support for specific workflows and information flows	High	High	Low	Low
Flexibility w.r.t. workflow and information flow	Low	Medium	High	High
User-level extensibility	Low	Medium	High	High
Ease of incorporation of outside code (code reuse)	Low	Low-Medium	Medium	High
Ease of experimentation	Low	Medium	Medium	High
Amount of new code required to create a complete simulation	Low	Medium	High	High (reuse can reduce)
Breadth of current "ecosystem" for "plugins"	Low	Medium	High	Low (but growing)
Ease of coupling simulations	Low	Low	Medium	High

26



Be Aware: “Framework” Describes Many Things

- Currently in scientific computing, this term means different things to different people
- **Basic software composition environment**
 - Examples: CCA, CORBA Component Model, ...
- An environment facilitating development of applications in **a particular scientific domain** (i.e. fusion, computational chemistry, ...)
 - Example: Earth System Modeling Framework, <http://www.esmf.ucar.edu>
 - Example: Computational Facility for Reacting Flow Science, <http://cfrfs.ca.sandia.gov>
- An environment for managing complex **workflows** needed to carry out calculations
 - Example: Kepler: <http://kepler-project.org>
- **Integrated data analysis and visualization environments** (IDAVEs)
- Lines are often fuzzy
 - Example: Cactus, <http://www.cactuscode.org>
- Others types of frameworks *could* be built based on a **basic software composition environment**



27



Relationships: Components, Objects, and Libraries

- Components are typically discussed as **objects** or collections of objects
 - **Interfaces** generally designed in **OO** terms, but...
 - Component **internals need not be OO**
 - **OO languages are not required**
- Component environments can **enforce** the use of **published interfaces** (prevent access to internals)
 - Libraries can not
- It is possible to load **several instances** (versions) of a component in a single application
 - Impossible with libraries
- Components *must* include some code to **interface with the framework/component environment**
 - Libraries and objects do not



28



What is the CCA?

- Component-based software engineering has been developed in other areas of computing
 - Especially business and internet
 - Examples: CORBA Component Model, COM, Enterprise JavaBeans
- Many of the needs are similar to those in HPC scientific computing
- But scientific computing imposes special requirements not common elsewhere
- CCA is a component environment specially designed to meet the needs of HPC scientific computing



Special Needs of Scientific HPC

- Support for legacy software
 - How much **change** required for component environment?
- Performance is important
 - What **overheads** are imposed by the component environment?
- Both parallel and distributed computing are important
 - What approaches does the component model support?
 - What **constraints** are imposed?
 - What are the **performance costs**?
- Support for **languages, data types, and platforms**
 - Fortran?
 - Complex numbers? Arrays? (as first-class objects)
 - Is it available on my parallel computer?






CCA: Concept and Practice

- In the following slides, we explain important **concepts** of component-based software from the CCA perspective
- We also sketch how these concepts are **manifested in code** (full details in the Hands-On)
- The **CCA Specification** is the mapping between concept and code
 - A standard established by the CCA Forum
 - Expressed in the Scientific Interface Definition Language (SIDL) for language neutrality (syntax similar to Java)
 - SIDL can be translated into bindings for specific programming languages using, e.g., the Babel language interoperability tool




CCA Concepts: Components



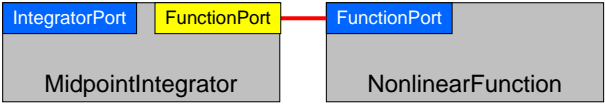
- A component encapsulates some computational functionality
- Components provide/use one or more interfaces
 - A component with no interfaces is formally okay, but isn't very interesting or useful
- In SIDL, a component is a **class** that **implements** (inherits from) **gov.cca.Component**
 - This means it must implement the **setServices** method to tell framework what ports this component will **provide** and **use**
 - **gov.cca.Component** is defined in the CCA specification





CCA
Common Component Architecture


CCA Concepts: Ports




```

graph LR
    MI[MidpointIntegrator] -- FunctionPort --- NF[NonlinearFunction]
    
```

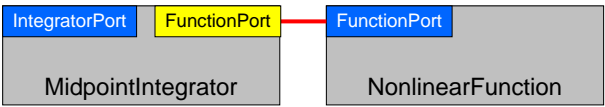
- Components interact through well-defined **interfaces**, or **ports**
 - A port expresses some **computational functionality**
 - In Fortran, a port is a bunch of subroutines or a **module**
 - In OO languages, a port is an **abstract class** or **interface**
- Ports and connections between them are a procedural (caller/callee) relationship, **not dataflow!**
 - e.g., *FunctionPort* could contain a method like `evaluate(in Arg, out Result)` with data flowing both ways


33



CCA
Common Component Architecture


CCA Concepts: *Provides* and *Uses* Ports

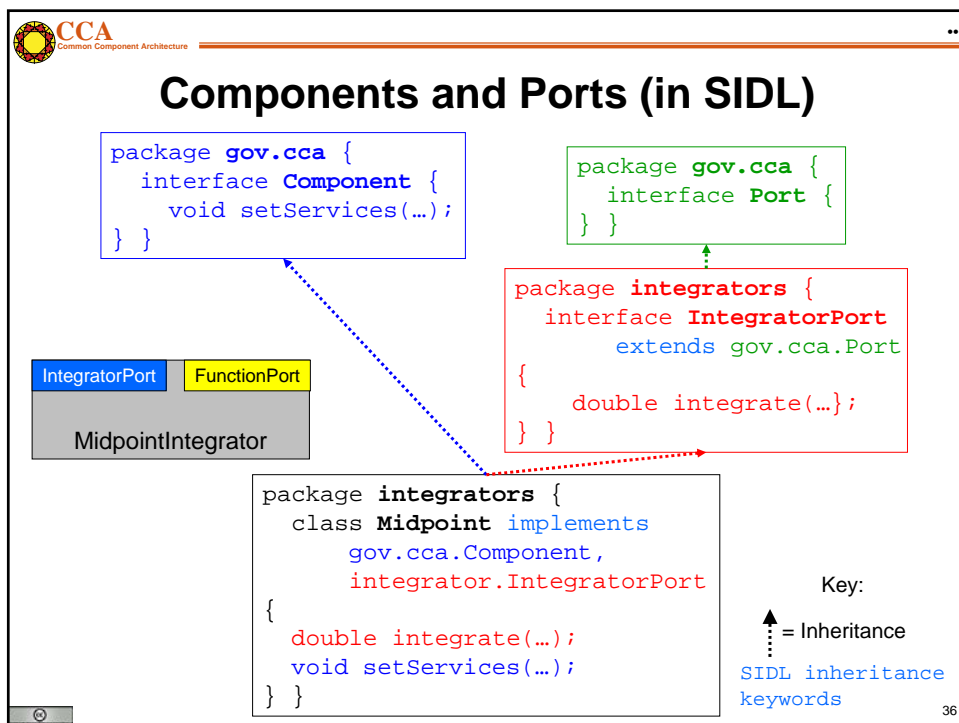
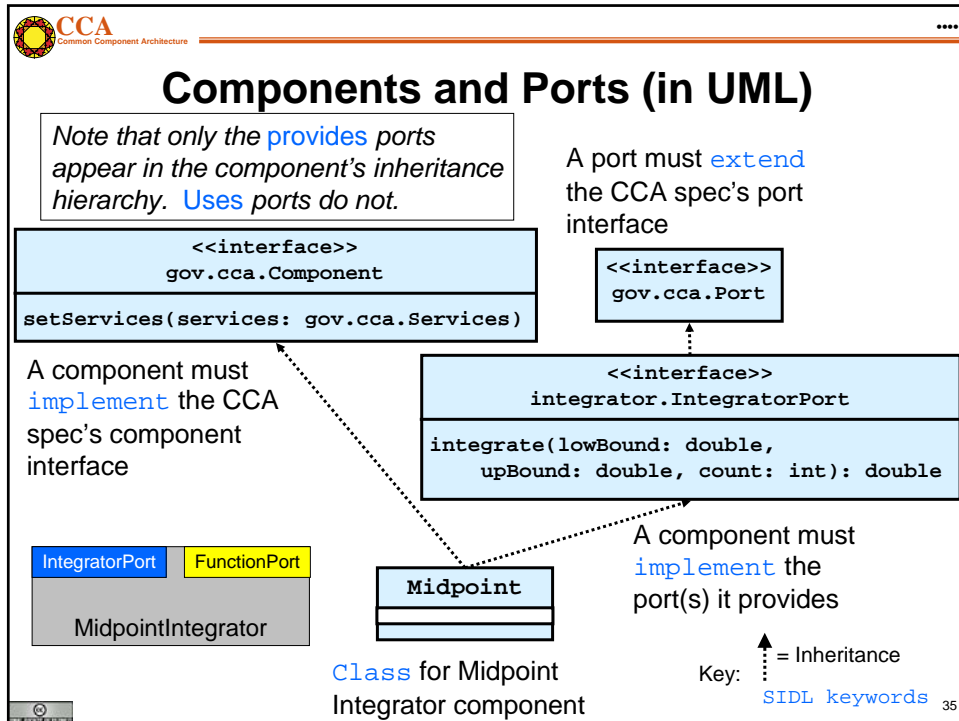



```

graph LR
    MI[MidpointIntegrator] -- FunctionPort --- NF[NonlinearFunction]
    
```

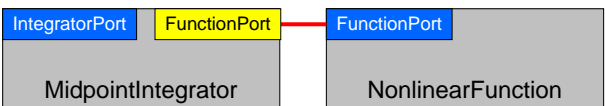
- Components may **provide** ports – **implement** the class or subroutines of the port (**"Provides" Port**)
 - *Providing* a port implies certain inheritance relationships between the component and the abstract definition of the interface (**more details shortly**)
 - A component can *provide* multiple ports
 - Different "views" of the same functionality, or
 - Related pieces of functionality
- Components may **use** ports – **call** methods or subroutines in the port (**"Uses" Port**)
 - *Use* of ports is just like calling a method normally except for a little additional work due to the "componentness" (**more details shortly**)
 - No inheritance relationship implied between caller and callee
 - A component can *use* multiple ports


34




 **CCA**
Common Component Architecture


Using Ports



```
graph LR; MI[MidpointIntegrator] -- IntegratorPort --- FP1[FunctionPort]; FP1 -- FunctionPort --- NF[NonlinearFunction];
```


- Calling methods on a port you use requires that you first obtain a “handle” for the port
 - Done by invoking `getPort()` on the user’s `gov.cca.Services` object
 - Free up handle by invoking `releasePort()` when done with port
- Best practice is to bracket actual port usage as closely as possible without using `getPort()`, `releasePort()` too frequently
 - Can be expensive operations, especially in distributed computing contexts
 - Performance is in tension with dynamism
 - can’t “re-wire” a ports that is “in use”

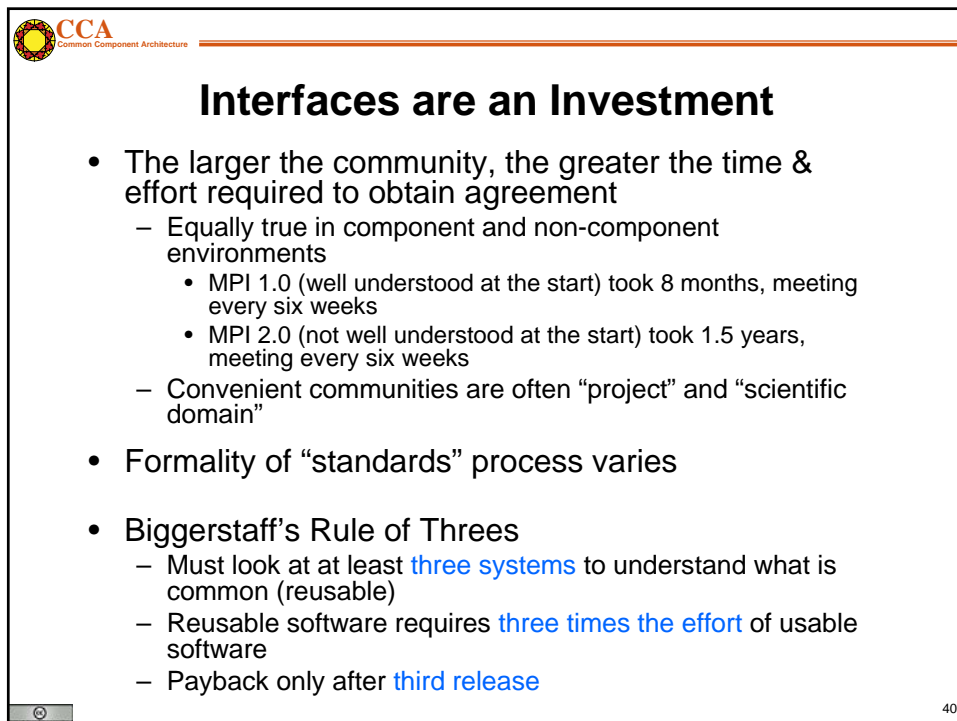
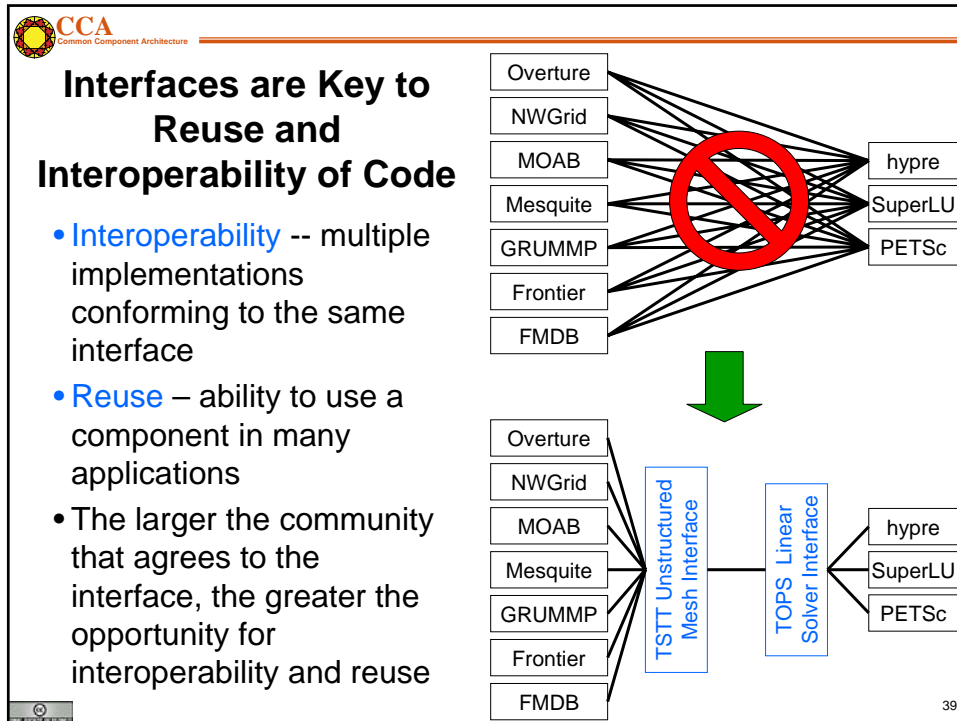
 37


 **CCA**
Common Component Architecture

Where Do Ports Come From?

- Most ports are designed and implemented by **users** of CCA
 - May be specific to an application or used more broadly (i.e. community-wide)
- The **CCA specification** defines a *small* number of ports
 - Most are services CCA frameworks must provide for use by components
 - Some are intended for users to implement in their components, and have a special meaning recognized by the framework
 - E.g. `gov.cca.ports.GoPort` provides a very simple protocol to start execution of component-based applications

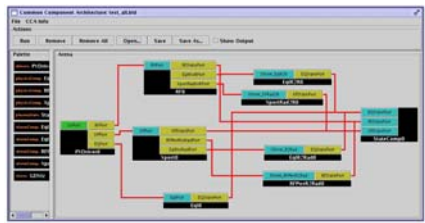
 38







CCA
Common Component Architecture

CCA Concepts: Frameworks



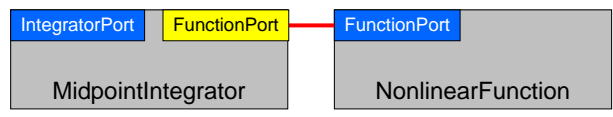
- The framework provides the means to “hold” components and **compose** them into applications
- Frameworks allow **connection of ports** without exposing component implementation details
- Frameworks provide a small set of **standard services** to components
 - Framework services are CCA ports, just like on components
 - Additional (non-standard) services can also be offered
 - Components can register ports as services using the *ServiceProvider* port
- *Currently:* specific frameworks are specialized for specific computing models (parallel, distributed, etc.)
- *Future:* better integration and interoperability of frameworks


41




CCA
Common Component Architecture

Components Must Keep Frameworks Informed




- Components must tell the framework about the ports they are providing and using
 - Framework will not allow connections to ports it isn't aware of
- Register them using methods on the component's **gov.cca.Services** object
 - **addProvidesPort()** and **removeProvidesPort()**
 - **registerUsesPort()** and **unregisterUsesPort()**
 - All are defined in the CCA specification
- Ports are usually registered in the component's **setServices()** method
 - Can also be added/removed dynamically during execution


42



CCA Concepts: Language Interoperability

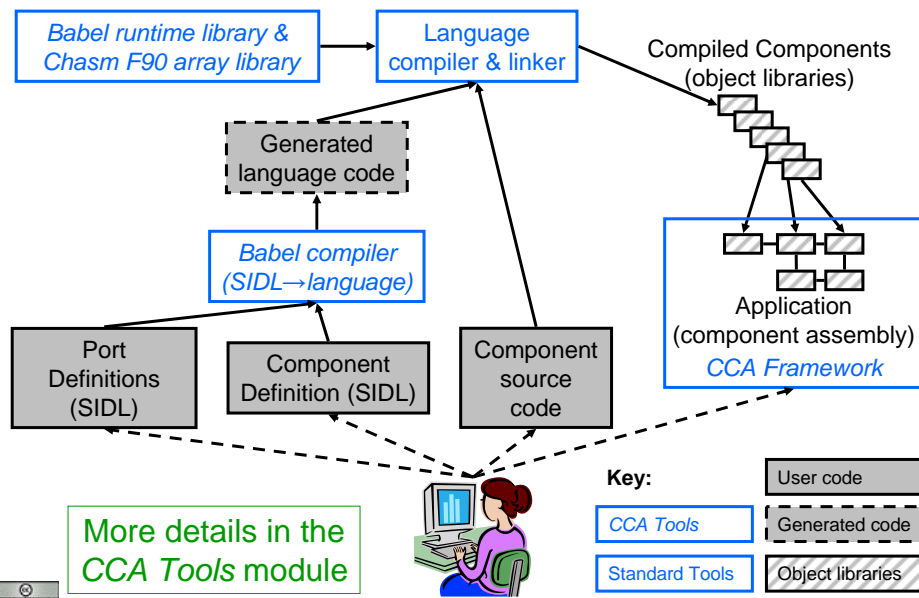
- Scientific software is increasingly diverse in use of programming languages
- In a component environment, users should not care what language a component is implemented in
- “Point-to-point” solutions to language interoperability are not suitable for a component environment
- The **Babel** language interoperability tool provides a common solution for all supported languages
- **Scientific Interface Definition Language** provides language-neutral way of expressing interfaces




43



Coding in a CCA Environment



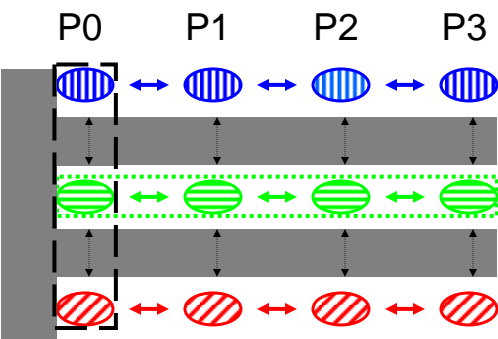
44



CCA Supports Parallelism -- by “Staying Out of the Way” of it

- Single component multiple data (SCMD) model is component analog of widely used SPMD model
- Each process loaded with the same set of components wired the same way
- Different components in same process “talk to each” other via ports and the framework


- **Same component in different processes talk to each other through their favorite communications layer (i.e. MPI, PVM, GA)**



Components: Blue, Green, Red
Framework: Gray

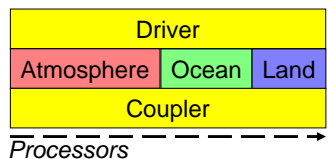
Any parallel programming environments that can be mixed outside of CCA can be mixed inside

45

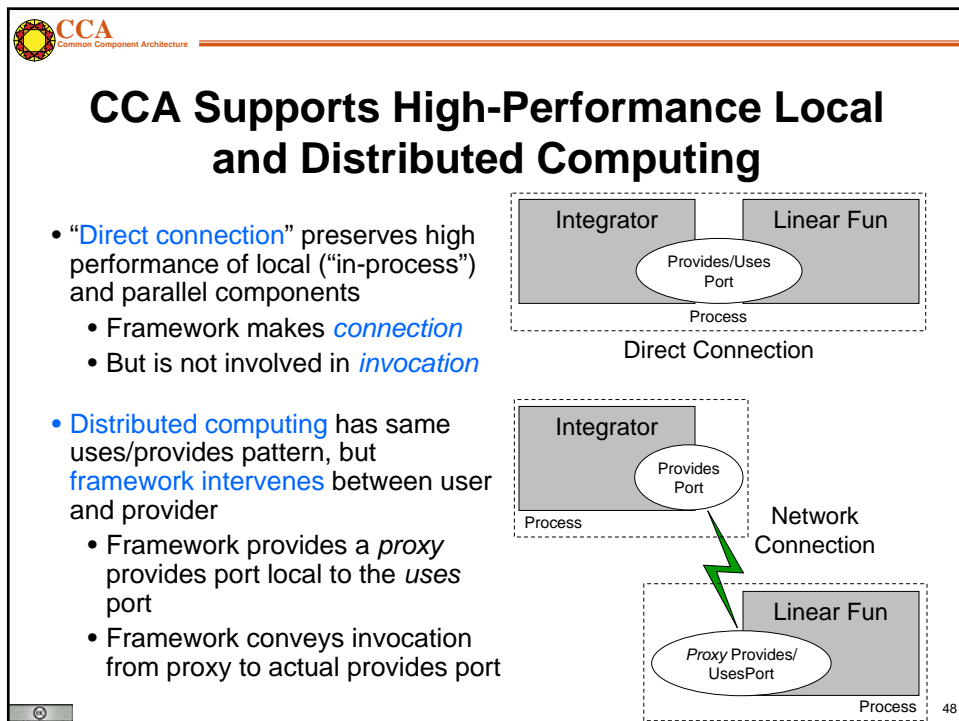
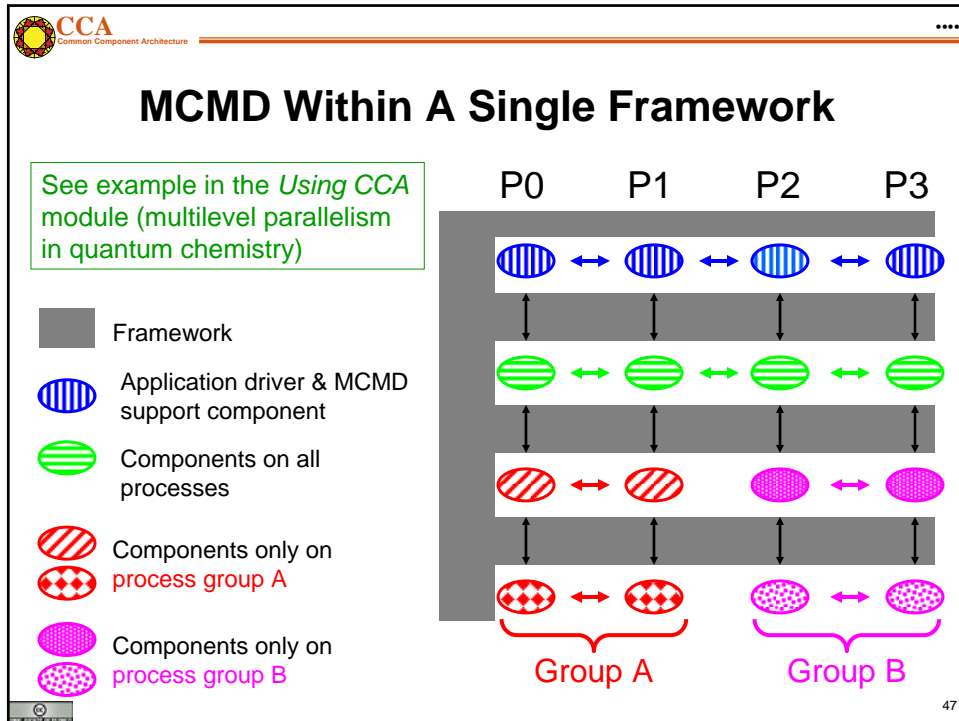


“Multiple-Component Multiple-Data” Applications in CCA

- Simulation composed of multiple SCMD sub-tasks
- Usage Scenarios:
 - Model coupling (e.g. Atmosphere/Ocean)
 - General multi-physics applications
 - Software licensing issues
 - i.e. limited number of instances
- Approaches
 - Run single parallel framework
 - Driver component that partitions processes and builds rest of application as appropriate (through BuilderService)
 - Run multiple parallel frameworks
 - Link through specialized communications components
 - Link as components (through AbstractFramework service)



46





“Direct Connection” Details

- Directly connected components are in the **same address space**
 - Data can be passed by reference instead of copying
 - Just like “traditional” programs
 - Framework involved in **connecting** components, but **not invocations** on ports
- Cost of “CCAness” in a direct connect environment is **one level of indirection** on calls **between** components
 - Equivalent to a C++ **virtual function call**: lookup function location, invoke it
 - Overhead is on the **invocation only** (i.e. latency), not the total execution time
 - Cost equivalent of ~2.8 F77 or C function calls
 - ~48 ns vs 17 ns on 500 MHz Pentium III Linux box



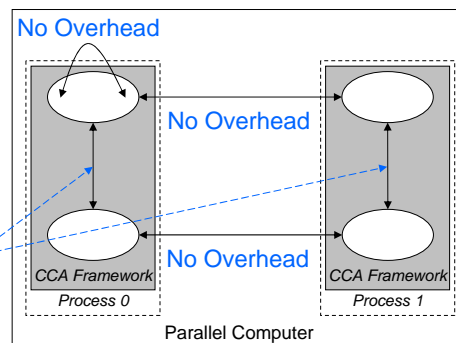
49



Performance, the Big Picture

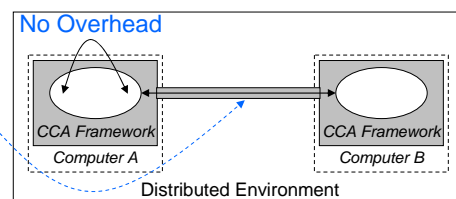
Direct-Connect, Parallel

- No CCA overhead on...
 - calls within component
 - parallel communications across components
- Small overheads on invocations on ports
 - Virtual function call (CCAness)
 - Language Interoperability (some data types)



Distributed

- No CCA overhead on calls within component
- Overheads on invocations on ports
 - Language interoperability (some data types)
 - Framework
 - (Wide area) network



50



Maintaining HPC Performance

- The performance of your application is as important to us as it is to you
- The CCA is designed to provide maximum performance
 - But the best we can do is to make your code perform **no worse**
- Facts:
 - Measured overheads per function call are **low**
 - Most overheads **easily amortized** by doing enough work per call
 - Other changes made during componentization may also have performance impacts
 - **Awareness** of costs of abstraction and language interoperability facilitates design for high performance

More about
performance in notes


51




Supplementary material for handouts

Some Performance Results and References

- Lois Curfman McInnes, et al. **Parallel PDE-Based Simulations Using the Common Component Architecture**. In Are Magnus Bruaset, Petter Bjørstad, and Aslak Tveito, editors, *Numerical Solution of PDEs on Parallel Computers*. Springer-Verlag, 2005. Invited chapter, in press.
- S. Benson, et al. **Using the GA and TAO Toolkits for Solving Large-Scale Optimization Problems on Parallel Computers**. Technical report ANL/MCS-P1084-0903, Argonne National Laboratory, September 2003.
- Boyana Norris, et al. **Parallel Components for PDEs and Optimization: Some Issues and Experiences**. *Parallel Computing*, 28:1811--1831, 2002.
- David E. Bernholdt, et al. **A Component Architecture for High-Performance Computing**. In *Proceedings of the Workshop on Performance Optimization via High-Level Languages and Libraries (POHLL-02)*, 2002.



Maximum 0.2% overhead for CCA vs native C++ code for parallel molecular dynamics up to 170 CPUs



Aggregate time for linear solver component in unconstrained minimization problem w/ PETSc

52



Overhead from Component Invocation

- Invoke a component with different arguments
 - Array
 - Complex
 - Double Complex
- Compare with f77 method invocation
- Environment
 - 500 MHz Pentium III
 - Linux 2.4.18
 - GCC 2.95.4-15
- Components took 3X longer
- Ensure granularity is appropriate!
- Paper by Bernholdt, Elwasif, Kohl and Epperly


Function arg type	f77	Component
Array	80 ns	224ns
Complex	75ns	209ns
Double complex	86ns	241ns

53



Scalability : Component versus Non-component. I

- Quantum chemistry simulation
- Sandia's MPQC code
 - Both componentized and non-componentized versions
- Componentized version used TAO's optimization algorithms
- Problem :Structure of isoprene HF/6-311G(2df,2pd)



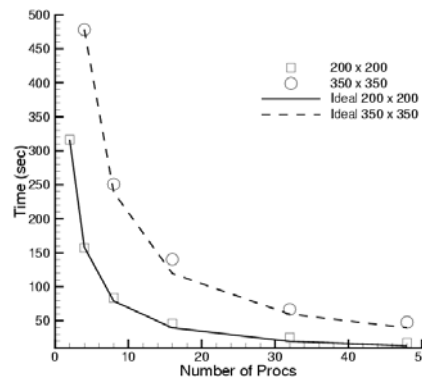
Parallel Scaling of MPQC w/ native and TAO optimizers

54



Scalability : Component versus Non-component. II

- Hydrodynamics; uses CFRFS set of components
- Uses GrACEComponent
- Shock-hydro code with no refinement
- 200 x 200 & 350 x 350 meshes
- Cplant cluster
 - 400 MHz EV5 Alphas
 - 1 Gb/s Myrinet
- Negligible component overhead
- Worst perf : 73% scaling efficiency for 200x200 mesh on 48 procs



Reference: S. Lefantzi, J. Ray, and H. Najm, Using the Common Component Architecture to Design High Performance Scientific Simulation Codes, *Proc of Int. Parallel and Distributed Processing Symposium*, Nice, France, 2003.



Advanced CCA Concepts

Brief introductions only, but more info is available – just ask us!

- Leveraging the component environment to provide additional capabilities to software developers
- The Proxy Component pattern (Hands-On Ch. 6, papers)
- Component lifecycle (tutorial notes, Hands-On)
- Components can be dynamic (papers)
- Improving the quality of component software (papers)
- Support for advanced parallel/high-performance computing (papers)



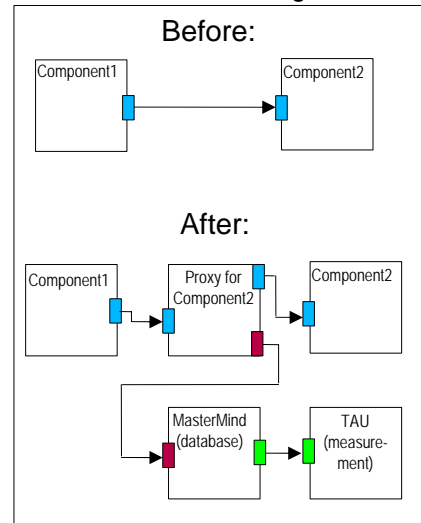
The Proxy Component Pattern

- A “proxy” component can be inserted between the user and provider of a port without either being aware of it (non-invasive)
- Proxy can **observe** or **act** on all invocations of the interface
- Similar to **aspect-oriented programming**
- For many purposes, proxies can be **generated automatically** from SIDL definition of the port

Sample uses for proxy components:

- **Performance:** instrumentation of method calls
- **Debugging:** execution tracing, watching data values
- **Testing:** Capture/replay

Performance Monitoring with TAU



57



Component Lifecycle

Additional
material
in notes

- **Composition Phase (assembling application)**
 - Component is **instantiated** in framework
 - Component interfaces are **connected** appropriately
- **Execution Phase (running application)**
 - Code in components uses functions provided by another component
- **Decomposition Phase (termination of application)**
 - **Connections** between component interfaces may be **broken**
 - Component may be **destroyed**

In an application, individual components may be in different phases at different times

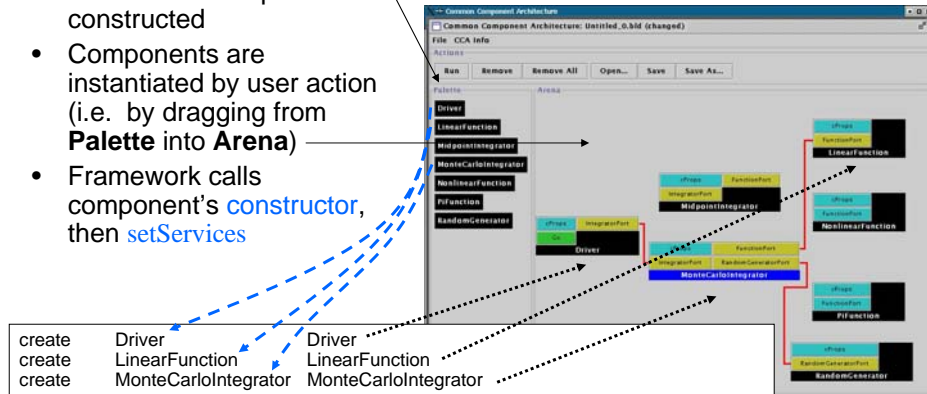
Steps may be under human or software control

58



User Viewpoint: Loading and Instantiating Components

- Components are code + metadata
- Using metadata, a **Palette** of available components is constructed
- Components are instantiated by user action (i.e. by dragging from **Palette** into **Arena**)
- Framework calls component's **constructor**, then **setServices**
- Details are **framework-specific!**
- **Ccaffeine** currently provides both command line and GUI approaches

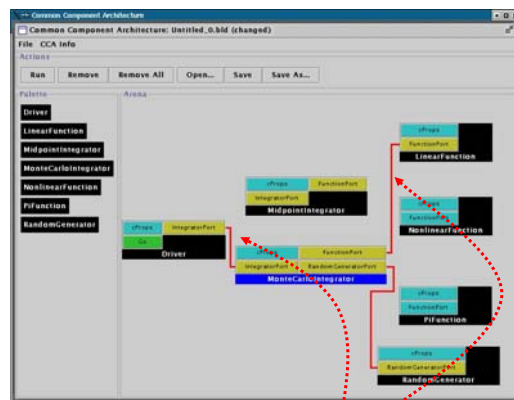


59



User Connects Ports

- Can only connect uses & provides
 - Not uses/uses or provides/provides
- Ports connected by type, not name
 - Port names must be unique within component
 - Types must match across components
- Framework puts info about *provider* of port into *using* component's Services object



connect	Driver	IntegratorPort	MonteCarloIntegrator	IntegratorPort
connect	MonteCarloIntegrator	FunctionPort	LinearFunction	FunctionPort
...				

60

CCA Common Component Architecture Supplementary material for handouts

Component's View of Instantiation

- Framework calls component's **constructor**
- Component initializes internal data, etc.
 - Knows *nothing* outside itself
- Framework calls component's **setServices**
 - Passes setServices an object representing everything "outside"
 - setServices declares ports component *uses* and *provides*
- Component *still* knows nothing outside itself
 - But Services object provides the means of communication w/ framework
- Framework now knows how to "decorate" component and how it might connect with others

61

CCA Common Component Architecture Supplementary material for handouts

Component's View of Connection

- Framework puts info about provider into **user component's** Services object
 - MonteCarloIntegrator's* Services object is aware of connection
 - NonlinearFunction* is not!
- MCI's** integrator code cannot yet call functions on FunctionPort

62

Supplementary material for handouts

Component's View of Using a Port

- User calls `getPort` to obtain (handle for) port from Services
 - Finally user code can “see” provider
- `Cast` port to expected type
 - OO programming concept
 - Insures type safety
 - Helps enforce declared interface
- `Call` methods on port
 - e.g.
 $\text{sum} = \text{sum} + \text{function} \rightarrow \text{evaluate}(x)$
- Call `releasePort`

The diagram illustrates the interaction between user code and the MonteCarloIntegrator. The user code calls `getPort` to obtain a port handle from the Services. This handle is then cast to the expected type. The user code then calls methods on the port, such as `function->evaluate(x)`, to perform calculations. Finally, the user code calls `releasePort` to release the port. The MonteCarloIntegrator component contains the Framework interaction code, the CCA.Services (which uses FunctionPort and is connected to NonlinearFunction FunctionPort), and the Integrator code.

63

Supplementary material for handouts

Dynamic Component Assemblies

- `gov.cca.BuilderService` allows **programmatic composition** of components
 - Components can be instantiated/destroyed, and connected/disconnected under program control

Sample uses of *BuilderService*:

- Python “**driver**” **script** which can assemble and control an application
 - i.e. MCMD climate model
- **Adaptation** to changing conditions
 - Swap components in and out to give better performance, numerical accuracy, convergence rates, etc.
 - TASCs project “**Computational Quality of Service**” activity

64



Enhancing Software Quality

- Current component architectures define **syntax** of interfaces
- Extend interface to include **semantics** (behavior) for more complete definition

– “Design by contract”

– Help ensure component **performs correctly**

– Help ensure component **is used correctly**

- Selective enforcement to control impact

- TASCs project “Software Quality and Verification” activity

```
package vector version 1.0 {
  class Utils { ...
    static double norm(in array<double> u,
                      in double tol,
                      in int badLevel)
    require /* Preconditions */
      not_null : u != null;
      u_is_1d : dimen(u) == 1;
      non_negative_tolerance : tol >= 0.0;
    ensure /* Postconditions */
      no_side_effects : is pure;
      non_negative_result : result >= 0.0;
      nearEqual(result, 0.0, tol)
      iff isZero(u, tol);
    ... }
}
```

65



Supporting Emerging HPC Hardware Environments

- CCA does not dictate a specific approach to parallelism
- Different approaches and tools can be provided via components and custom frameworks

Examples...

- Uintah Computational Framework (Utah) provides a multi-threaded parallel execution environment based on task graphs
 - Specialized to certain structured adaptive mesh refinement problems
- TASCs developing services to manage groups of parallel components/tasks (MCMD)

Also...

- TASCs developing support for heterogeneous processor environments
 - FPGAs, GP-GPUs, accelerators, and other co-processors
 - Accelerator code encapsulated as components, interacting w/ components on primary processors
- Integration of fault tolerance capabilities with CCA under development (CIFTS-TASCs collaboration)

66



Is CCA for You?

- Much of what CCA does can be done without such tools *if* you have sufficient discipline
 - The larger a group, the harder it becomes to impose the necessary discipline
- Projects may use different aspects of the CCA
 - CCA is *not* monolithic – use what *you* need
 - Few projects use all features of the CCA... initially
- Evaluate what *your* project needs against CCA's capabilities
 - Other groups' criteria probably differ from yours
 - CCA continues to evolve, so earlier evaluations may be out of date
- Evaluate CCA against other ways of obtaining the desired capabilities
- Suggested starting point:
 - CCA tutorial "hands-on" exercises

67



Take an Evolutionary Approach

- The CCA is designed to allow selective use and incremental adoption
- "SIDLize" interfaces incrementally
 - Start with essential interfaces
 - Remember, only externally exposed interfaces need to be Babelized
- Componentize at successively finer granularities
 - Start with whole application as one component
 - Basic feel for components without "ripping apart" your app.
 - Subdivide into finer-grain components as appropriate
 - Code reuse opportunities
 - Plans for code evolution

68



View it as an Investment

- CCA is a long-term investment in your software
 - Like most software engineering approaches
- There is a cost to adopt
- The payback is longer term
- Remember Biggerstaff's Rule of Threes
 - Look at three systems, requires three times the effort, payback after third release



69



CCA is Still Under Development

- We've got...
 - A stable component model
 - Working tools
 - Active users
- But...
 - We know its not perfect
 - We're not "done" by any stretch
- Talk to us...
 - If you're evaluating CCA and and need help or have questions
 - If you don't think CCA meets your needs
 - If you've got suggestions for things we might do better



70



What Can CCA Do Today?

- Ccaffeine framework for HPC/parallel
 - XCAT and other options for distributed computing
- Language interoperability
 - Fortran 77/90/95, C, C++, Java, Python
 - Support for Fortran/C user-defined data structures under development
- CCA Tools working on a variety of platforms
 - Linux most widely used
 - Mac OS X second
 - Some IBM AIX users
 - Works on Cray XT series, port to IBM BlueGene in progress
 - Porting is driven by user needs, so let us know!



71



CCA
Common Component Architecture

CCA Tools – Language Interoperability and Frameworks

CCA Forum Tutorial Working Group

<http://www.cca-forum.org/tutorials/>
tutorial-wg@cca-forum.org



This work is licensed under a [Creative Commons Attribution 2.5 License](http://creativecommons.org/licenses/by/2.5/)

72

CCA
Common Component Architecture

Goal of This Module

- Describe **tools** for multi-lingual, scientific component 'plug-and-play'

IDE = Interactive Development Environment

73

CCA
Common Component Architecture

CCA adds value to component development

Babel runtime library & Chasm F90 array library

Language compiler & linker

Compiled Components (object libraries)

Generated language code

Babel compiler (SIDL → language)

Port Definitions (SIDL)

Component Definition (SIDL)

Component source code


Application (component assembly)
CCA Framework

CCA IDE

Key:

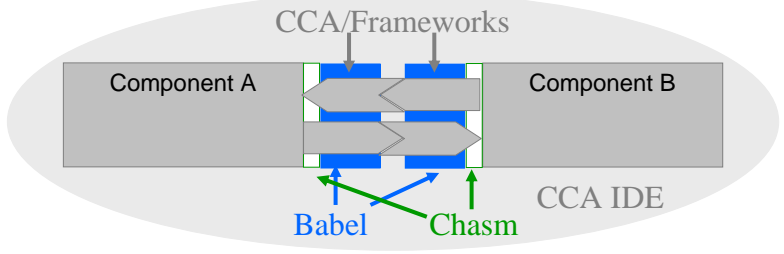
- User code
- CCA Tools
- Generated code
- Standard Tools
- Object libraries

74

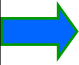


CCA
Common Component Architecture


Tools Module Overview



- Language interoperability tools
- Frameworks
- CCA Interactive Development Environment



75

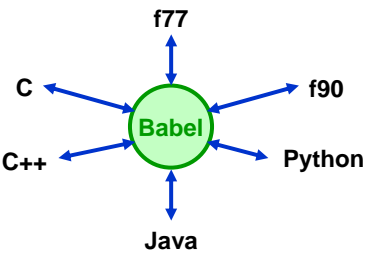


CCA
Common Component Architecture


Babel

Babel Facilitates Scientific Programming Language Interoperability


- Programming language-neutral interface descriptions
- Native support for basic scientific data types
 - Complex numbers
 - Multi-dimensional, multi-strided arrays
- Automatic object-oriented wrapper generation



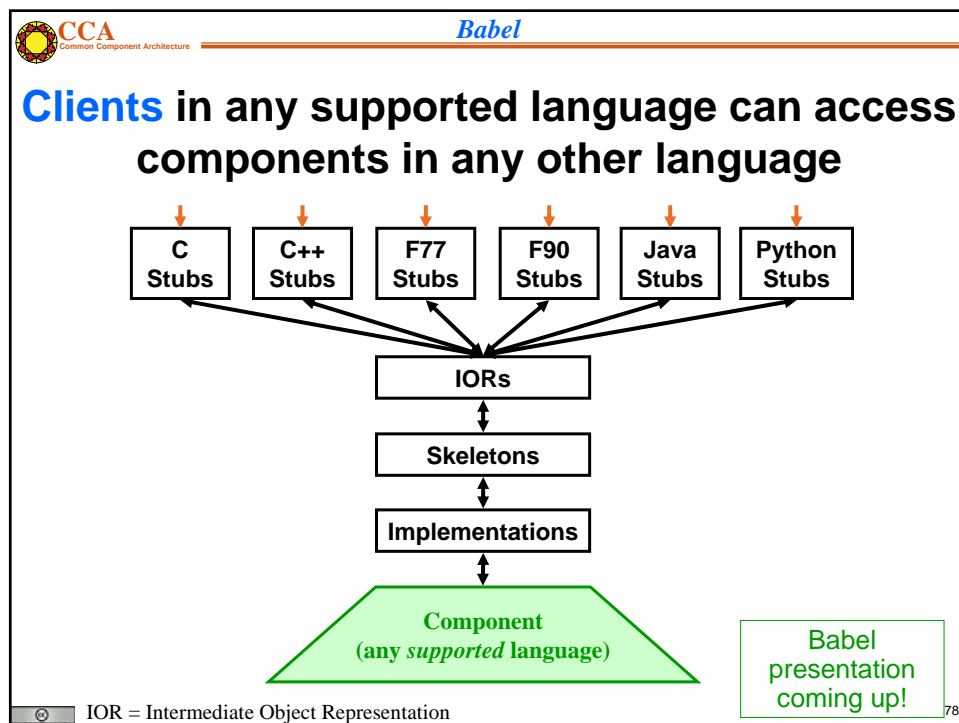
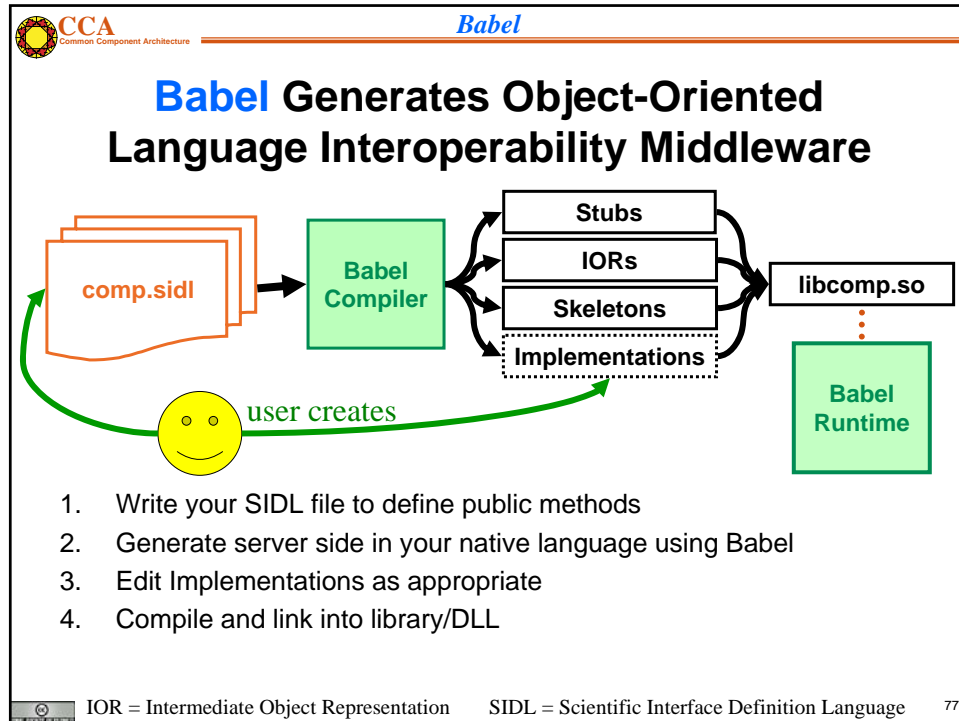
vs.



Supported on Linux, AIX, and Solaris 2.7, works on OSX;
 C (ANSI C), C++ (GCC), F77 (g77, Sun f77), F90 (Intel, Lahey, GNU, Absoft), Java (1.4)



76






Recent and Upcoming Features

- Remote Method Invocation (Babel RMI)
Kumfert, Leek, & Epperly IPDPS'07
ALPHA – Babel 1.0. (aka “stable branch”)*
*BETA – Babel 1.1.**
Final – Future Stable Branch (Babel 1.2?)
- Design-by-Contract
Dahlgren CBSE'07 (to appear)
ALPHA – Internal Babel branch
BETA – Babel 1.1.2 (expected)
- Feature requests are accepted and tracked in the open using Roundup
<http://www.cca-forum.org/bugs/babel>




Tools Module Overview



- Language interoperability tools
- ➔ • Frameworks
- CCA Interactive Development Environment






CCA
Common Component Architecture

Frameworks are Specialized to Different Computing Environments


- “**Direct connection**” preserves high performance of local (“in-process”) and parallel components
 - Framework makes **connection**
 - Framework *not* involved in **invocation**
- **Distributed computing** has same uses/provides pattern, but the framework **intervenes** between user and provider
 - Framework provides a **proxy port** local to the user's **uses** port
 - Framework conveys invocation from **proxy** to **actual** **provides** port



Direct Connection

Network Connection

81




CCA
Common Component Architecture

Graphical User Interfaces (GUIs) Deliver Plug-and-Play Experience

- Plug & play for:
 - Application software assembly
 - Visualization pipelines
 - Workflow management
- Assembling “wiring” diagrams is almost universal.
 - Software assembly: Ccaffeine, XCAT, SciRUN
 - Workflow: XCAT, SciRUN
 - Visualization: SciRUN

None of these (yet) plug into your favorite Integrated Development Environment (e.g., Eclipse, MS Dev. Studio, Java Studio, ...).

82


 **CCA**
Common Component Architecture


Caffeine Framework

Caffeine is a *Direct-Connect*, Parallel-Friendly Framework

- Supports SIDL/Babel components
 - Conforms to latest CCA specification (0.7)
 - Also supports legacy CCA specification (0.5)
 - Any C++ allowed with C and Fortran by C++ wrappers
- Provides command-line and GUI for composition
 - Scripting supports batch mode for SPMD
 - MPMD/SPMD custom drivers in any Babel language

Supported on Linux, AIX, OSX and is portable to modern UNIXes.

 83


 **CCA**
Common Component Architecture


Caffeine Framework

Caffeine Involves a Minimum of Three Steps

- As easy as 1-2-3:
 - Write your component.
 - Describe it with an XML file.
 - Add a line to your Caffeine input file to load the component at runtime.
- Proceed to plug & play...

There are many details and automated tools that help manage components.

 84



CCA
Common Component Architecture

Ccaffeine Framework (ccafe-gui)


Optional Ccaffeine GUI

- Process
 - User input is broadcast SPMD-wise from Java.
 - Changes occur in GUI *after* the C++ framework replies.
 - If your components are computing, GUI changes are blocked.
- Components interact through *port connections*
 - *provide* ports *implement* class or subroutines "Provides" Port
 - *use* ports *call* methods or subroutines in the port. "Uses" Port
 - Links denote caller/callee relationship *not* data flow

IntegratorPort

FunctionPort

MidpointIntegrator




FunctionPort

NonlinearFunction

Java is required.

85



CCA
Common Component Architecture

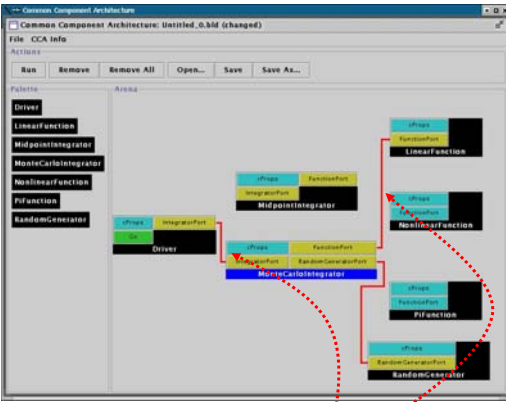
Ccaffeine Framework (ccafe-gui)

Supplementary material for notes

User Connects Ports

- Can only connect *uses* & *provides*
 - *Not* uses/uses
 - *Not* provides/provides
- Ports connected by type not name
 - Port names *must be unique* within component
 - Types *must match* across components
- Framework puts info about *provider* of port into *using* component's Services object

connect	Driver	IntegratorPort	MonteCarloIntegrator	IntegratorPort
connect	MonteCarloIntegrator	FunctionPort	LinearFunction	FunctionPort
...				



86

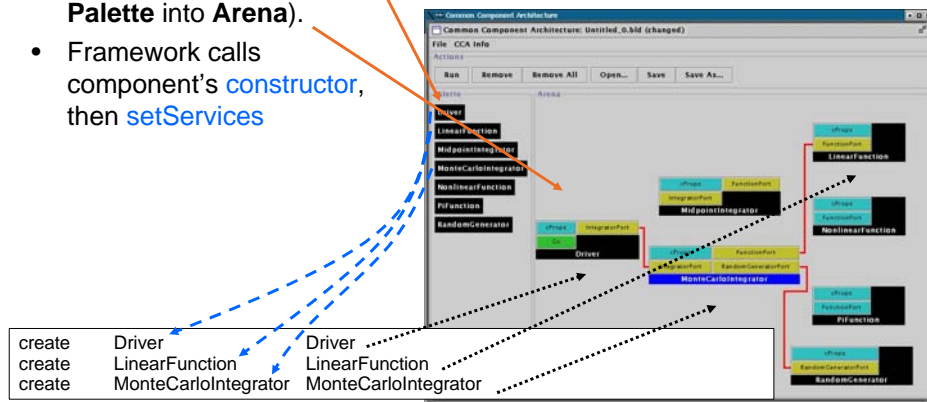


Common Component Architecture

Ccaffeine Framework (ccafe-gui)

Building an Application (1 of 2)

- Components are code + XML metadata
- Using metadata, a **Palette** of available components is constructed.
- Components are instantiated by user action (i.e. by dragging from **Palette** into **Arena**).
- Framework calls component's **constructor**, then **setServices**



87



Common Component Architecture

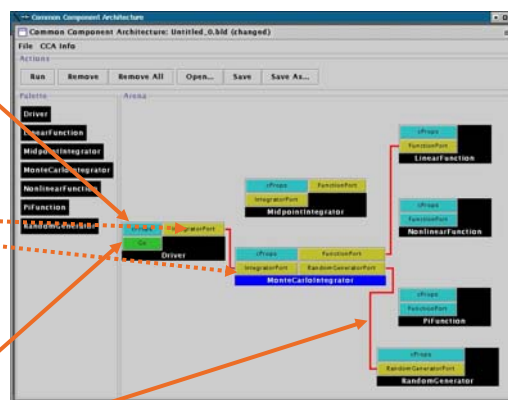
Ccaffeine Framework (ccafe-gui)

Building an Application (2 of 2)

1. Click **Configure** port to start parameter input dialogue.


2. For each connection:
click a **uses** port
then click a **provides** port
to establish a connection.

3. Click **Go** port to start the application.




Right-clicking a connection line breaks the connection -- enabling component substitution.

88

**CCA**
Common Component Architecture

XCAT-C++ Framework



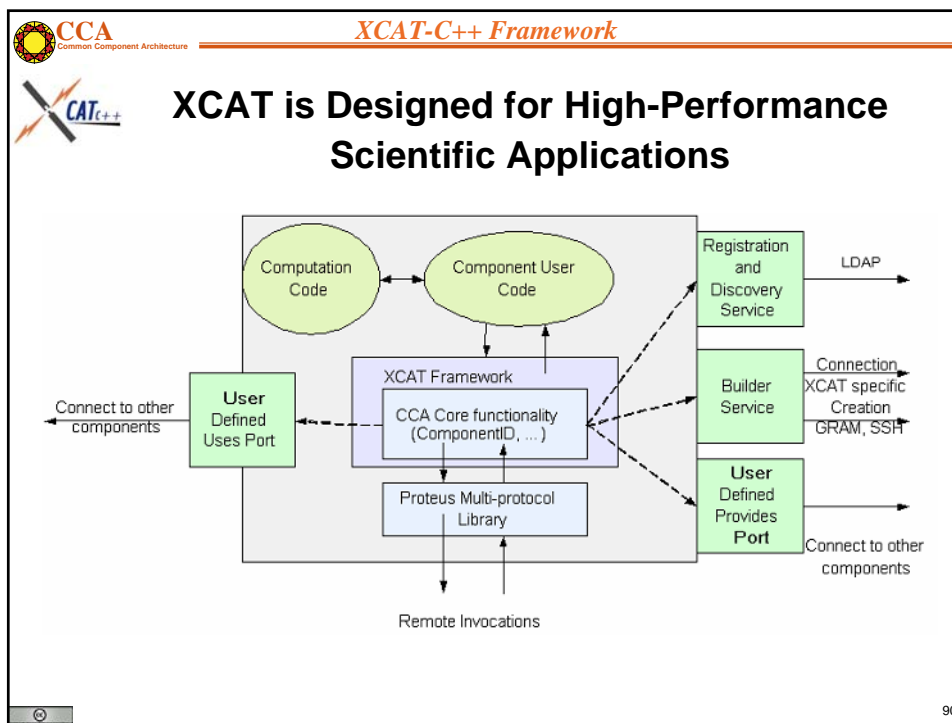
XCAT is a Web-services based Distributed Component Framework


- Remote references
 - Port types described in C++ header files or in WSDL documents
- User Interface
 - C++ and Python interface to CCA *BuilderService*
 - Uses SWIG for Python-C++ translations
- Component creation
 - Remote creation via SSH
- Component communication
 - Proteus multi-protocol library
 - Communication libraries can be loaded at run-time

Tested on Linux.

WSDL = Web Service Definition Language


89



 **CCA**
Common Component Architecture

XCAT-C++ Framework


 91

 **CCA**
Common Component Architecture

XCAT-C++ Framework

Supplementary material for notes

 92




CCA
Common Component Architecture

SCIRun2 Framework

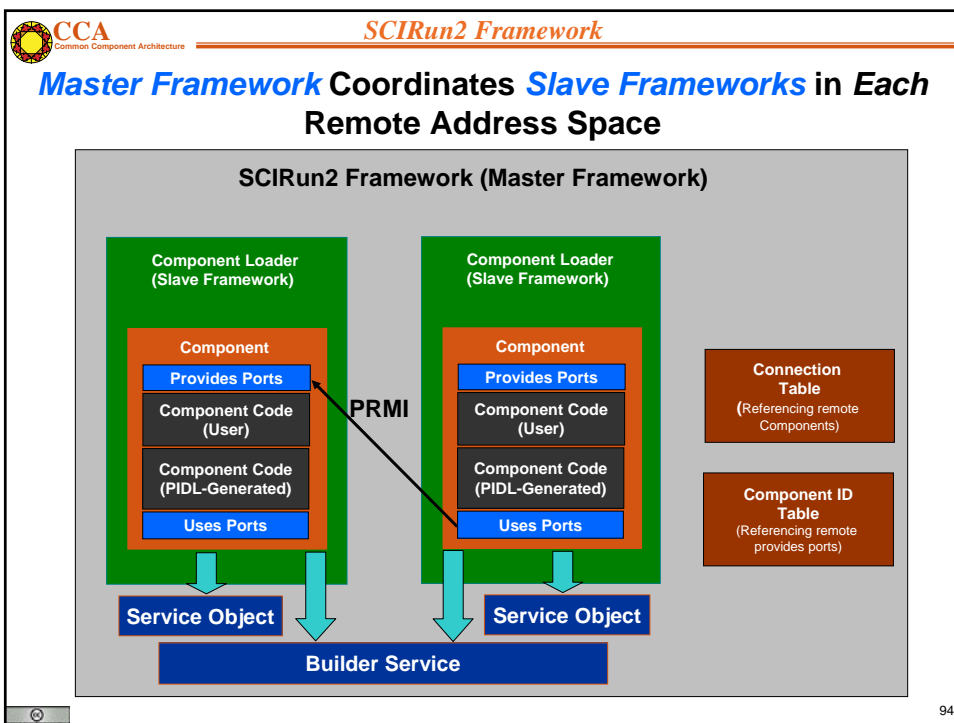
SCIRun2 is a Cross-Component Model, Distributed Component Framework


- Semi-automatically bridges component models
 - Templated components connected at run-time generate bridges
- Parallel Interface Definition Language (PIDL) – a SIDL variant
- User interface – GUI and textual
 - Dynamic composition
- Component and framework creation
 - Remote via SSH
- Component communication
 - C++ RMI with co-location optimization
 - MPI/ Parallel Remote Method Invocation (PRMI)



Supported on Linux.

93





 **CCA**
Common Component Architecture

SCIRun2 Framework


Basic How-To

1. Add component source files and makefile to SCIRun2 sources
 - May need to define ports in SIDL
2. Add component information to the component model xml file
3. Build component using SCIRun2 make scripts
 - Alternatively, build component using Babel
4. Start the framework and graphical (default) or text builder
5. Graphically connect component to other CCA-based or non CCA-based components
 - May need to create bridge components to go between models
6. Press the “Go” button on the driver component

 95

 **CCA**
Common Component Architecture


Tools Module Overview




The diagram illustrates the Tools Module Overview. It features a large light blue oval containing two gray rectangular blocks labeled 'Component A' and 'Component B'. Between these components are several interlocking arrow shapes representing interfaces. Above the oval, the text 'CCA/Frameworks' has two arrows pointing down to the interface area. Below the oval, the text 'CCA IDE' is in purple. Two arrows labeled 'Babel' and 'Chasm' point upwards from below the oval towards the interface area between the components.

- Language interoperability tools
- Frameworks

➡ • CCA Development Environment

 96

**CCA**
Common Component Architecture


CCA IDE

Component Development Environment Provided via **Eclipse** Plug-ins

- Provides a high-level graphical environment
 - Creating new SIDL-based components
 - Componentizing legacy codes
 - C, C++, Java and Fortran
- Automatic code generation

Supported on Linux, Windows, MacOS.

97

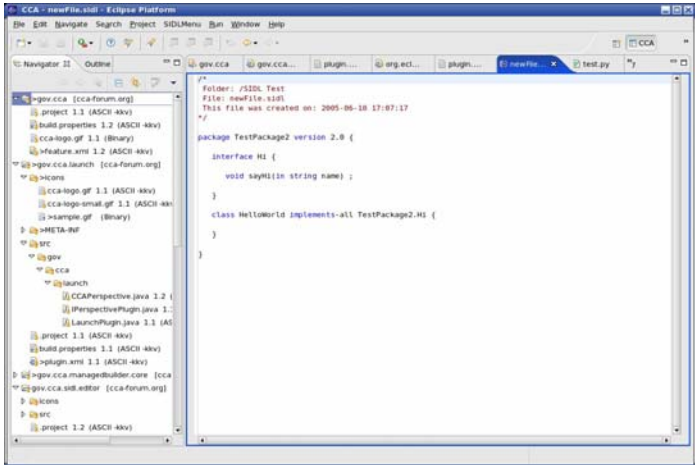
**CCA**
Common Component Architecture

CCA IDE

Component Development Environment Starts at the Eclipse Platform Level

Plug-ins for:

- SIDL Editor
- Wizards
- Preliminary automated build support



```
/*
 * Folder: /SIDL Test
 * File: newFile.sidl
 * This file was created on: 2005-06-18 17:07:17
 */
package TestPackage2 version 2.0 {
    interface HI {
        void sayHi(in string name);
    }
    <class HelloWorld implements all TestPackage2.HI {
    }
}
```


Imperative that you start by creating a new project!

98



Wizards are Available for Adding Packages and Classes or Generating SIDL from Legacy Codes

- Intuitive interfaces to port and component definition
- Helper wizards for setting port, component and (in the future) application properties



99



Bocca Creates “Scaffolding” for an entire CCA application

- Including **ports**
 - Give the typename and an empty port is created
- Including **components**
 - Give the name and an empty component is created
 - With extra parameters the component uses/provides a port
- Including **build system**
 - For all ports/components in the project
 - Implemented in any CCA supported language
- Including **application composition**
 - Link components to create application with Ccaffeine GUI

100



Bocca is an Application Generator

- Purely command-line driven
 - Generates the glue code that enables components and language interoperability
 - Enables rapid generation of an entire application and build system in minutes
 - Good for rapid prototyping and novice users...
 - But is meant to carry through to full project life-cycle
- Different and complementary to an IDE
 - Does not even have an editor for creating code (choose your own)




101



Bocca Demo




102


 **CCA**
Common Component Architecture

Supplementary material for notes

CCA Tools Contacts (1 of 2)

Tool	Purpose	More information
Babel	Scientific language interoperability tool kit	URL: www.llnl.gov/CASC/components Email: components@llnl.gov or babel-users@lists.llnl.gov
Ccaffeine	Direct-connect, parallel-friendly framework	URL: www.cca-forum.org/software/ Email: Ben Allan, ccafe-help@z.ca.sandia.gov Wiki: https://www.cca-forum.org/wiki
CCA IDE	CCA development environment	Email: usability@cca-forum.org


 103


 **CCA**
Common Component Architecture

Supplementary material for notes


CCA Tools Contacts (2 of 2)

Tool	Purpose	More information
SCIRun2	Cross-component model framework	URL: www.sci.utah.edu/ Email: Steve Parker, sparker@cs.utah.edu
XCAT-C++	Globus-based GRID framework	URL: grid.cs.binghamton.edu/projects/xcat/ Email: Madhu Govindaraju, mgovinda@cs.binghamton.edu
Bocca	Development environment	Email: bocca-dev@cca-forum.org

 104


 **CCA**
Common Component Architecture

Module Summary




The diagram illustrates the Module Summary architecture. It features two green rectangular blocks labeled 'Component A' and 'Component B' connected by three orange arrows pointing in both directions, representing the 'CCA/Frameworks'. Below these components, two blue arrows labeled 'Babel' and 'Chasm' point towards the connection. The entire setup is enclosed in a light purple oval labeled 'CCA IDE'.


- Described **tools** for multi-lingual, scientific component 'plug-and-play'
 - Language interoperability through Babel and Chasm
 - CCA Frameworks provide mechanisms for composition
 - CCA Interactive Development Environment via Eclipse plug-in




105

 **CCA**
Common Component Architecture

Boneyard Slides



106




CCA
Common Component Architecture


Supplementary material for notes

Experimental Frameworks

Framework	Purpose	Summary
Distributed CCA Framework (DCA)	MxN research	<ul style="list-style-type: none"> • Goal: explore MxN Parallel-Remote Method Invocation (PRMI) using MPI • Parallel data transfer and redistribution integrated into port invocation mechanism
LegionCCA	Grid-based research	<ul style="list-style-type: none"> • Goal: allow component-based CCA applications to run in Grid-scale environments using Legion • Supports creation, scheduling, persistence, migration, and fault notification; relies on Legion's built-in RPC mechanism (~Unix sockets)
XCAT-Java	Globus-based Grid research	<ul style="list-style-type: none"> • Goal: explore web interface for launching distributed applications • This (alpha) version compatible with latest CCA specification and provides built-in seamless compatibility with OGSi.


OGSI = Open Grid Services Infrastructure

107



CCA
Common Component Architecture

Chasm


Supplementary material for notes

Chasm-Assisted Glue Code Generation


1. Create functions_LinearFunction.impl.xml
2. Create xml description of source procedures


```
% gfortran -fdump-parse-tree LinearFunction.f90 >
LinearFunction.ptd
% gfortran2xml < LinearFunction.ptd >
LinearFunction.xml
```
3. Create .sidl, _Impl.F90, and _Mod.F90


```
% xalan -o functions_LinearFunction.sidl
functions_LinearFunction.impl.xml
cca-f90-comp.sidl.xsl
% xalan -o functions_LinearFunction_Impl.F90
functions_LinearFunction.impl.xml
cca-f90.impl.xsl
% xalan -o functions_LinearFunction_Mod.F90
functions_LinearFunction.impl.xml
cca-f90.mod.xsl
```
4. Run Babel...



108


CCA
Common Component Architecture

Chasm

Supplementary material for notes

User-Created XML Component Description File

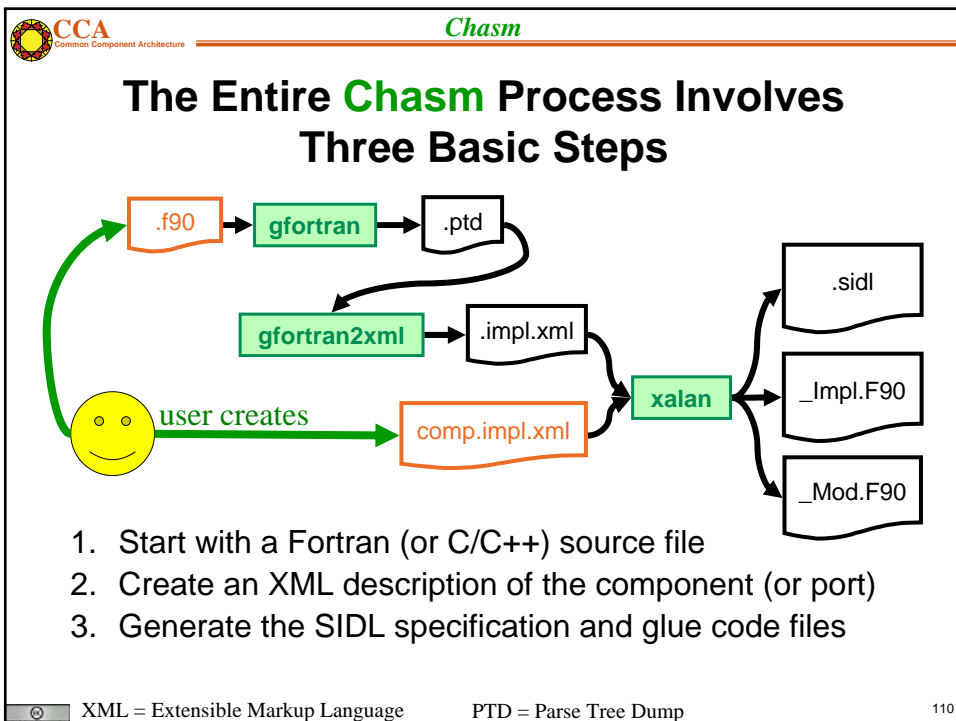
```

<componentImplementation name="LinearFunction" package="functions">
  <language name="F90">
    <property name="impl-scope" value="LinearFunction"/>
    <property name="impl-xml" value="/home/cca/LinearFunction.xml"/>

    <ports>
      <provides name="FunctionPort" package="function">
        <MethodsBlock>
          <Method name="evaluate" impl="evaluate_if"/>
        </MethodsBlock>
      </provides>
    </ports>

  </language>
</componentImplementation>
  
```

109





Chasm Provides Language Interoperability for Fortran, C, and C++

- Extracts interfaces from C/C++ and Fortran90 codes
- Uses library of XSLT stylesheets for language transformations → easily extended
 - Generates XML and SIDL representations
 - Generates Fortran90 Babel implementation glue
- Provides Fortran array descriptor library used by Babel

Supported on Linux, AIX, and Solaris 2.7, works on OSX;
C (ANSI C), C++ (GCC), F90 (Intel, Lahey, GNU, Absoft)



A Wizard is also Available for Adding Methods

SIDL File Wizard
This page is for adding a method

Method Properties
Method Name:
Return Type:

If your argument type is an array, you need to specify the following:
Array Type: Dimension: Order:


Arguments
Argument Name:
Argument Type: Passing Mode:

If your argument type is an array, you need to specify the following:
Array Type: Dimension: Order: Your current arguments:

Exceptions
Here is a list of exceptions. Please choose the one(s) your method throws, if any:
Exceptions: ☐ Exception0 ☐ Exception1 ☐ Exception2 ☐ Exception3 ☐ Exception4

☐ I would like to add another method.

What would you like to add next?
Add a Package:


 **CCA**
Common Component Architecture

CCA IDE


Supplementary material for notes

Recent and Upcoming Features

- Provide automated build support *1H 2005*
- Launch application via GUI *1H 2006*




113

 **CCA**
Common Component Architecture

SCIRun2 Framework

Supplementary material for notes


Simple SCIRun2 CCA (PIDL) and Babel Bridge




The screenshot shows the SCIRun2 IDE interface. The top menu bar includes 'File', 'Proxy Frameworks', 'CCA', 'Corba', 'Dataflow', 'Tao', 'Vtk', and 'babel'. A 'Messages' window on the right displays the following text:

```
SCIRun2  
Framework URL: socket://ottawa.sci.utah.edu:51484/134768216
```

The main workspace contains a component diagram with three components: 'Hello Component', 'Hello_NewPort', and 'NewPortConn'. 'Hello Component' is connected to 'Hello_NewPort', which is connected to 'NewPortConn'. A 'Hello' dialog box is open, showing the message 'NewPort: StringPort_impl: getString CALLED' and an 'OK' button.



114




CCA
Common Component Architecture

Ccaffeine Framework (ccafe-gui)


Supplementary material for notes


Application Configurations can be Re-used

1. Click **Save** or **Save As...** to save actions.
2. Click **Open** to replay actions.




- Script optimization
`% simplify-bld saved_file.bld > faster_file.bld`
- Batch conversion
`% bld2rc faster_file.bld > faster_file.batch`
- C++ stand-alone execution
`% bld2babel-cpp faster_file.bld faster_file_babel outdir`
or `% bld2neo faster_file.bld faster_file.batch outdir`


115




CCA
Common Component Architecture



CCA Forum Tutorial Working Group

<http://www.cca-forum.org/tutorials/>
tutorial-wg@cca-forum.org


This work is licensed under a [Creative Commons Attribution 2.5 License](https://creativecommons.org/licenses/by/2.5/)
116



Goal of This Module

Expose Attendees to How Existing Code is

- Wrapped into Babel objects, and then
- Promoted to CCA components

In the process, will also cover

- Babel Basics
- Scientific Interface Definition Language (SIDL)
- Code sample (“Hello World!”)



117




BABEL is Middleware for HPC



2006

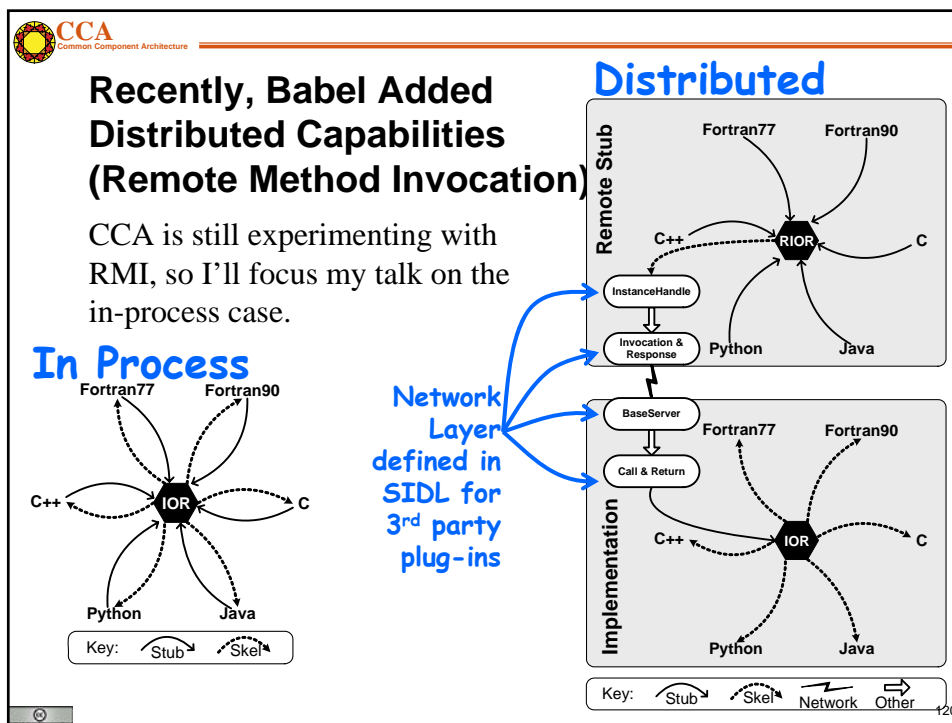
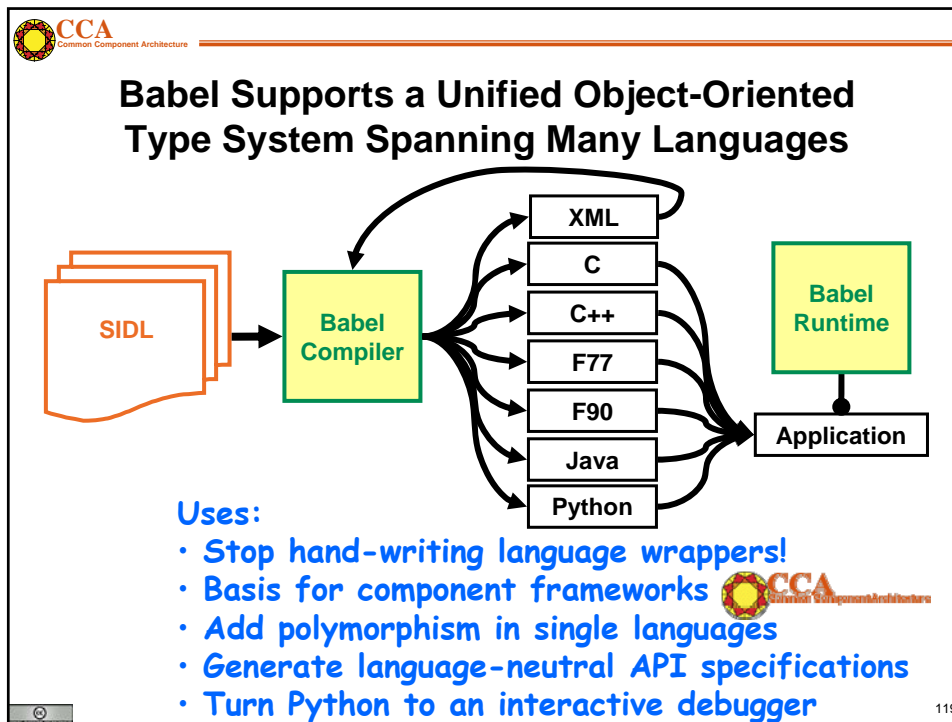
“The world’s most rapid communication among many programming languages in a single application.”




	CORBA	COM	.NET	Babel
BlueGene, Cray, Linux, AIX, & OS X	No	No	No	Yes*
Fortran	No	Limited	Limited	Yes
Multi-Dim Arrays	No	No	No	Yes
Complex Numbers	No	No	No	Yes
Licensing	Vendor Specific	Closed Source	Closed Source	Open Source

* BlueGene port expected mid 2007¹⁸







CCA
Common Component Architecture

Working Code: “Hello World” in F90 Using a Babel Type

```

program helloclient
  ① use greetings_english
  ② use sihl_BabelInterface
  implicit none
  ③ type(greetings_english_t) :: obj
  ④ type(sihl_BabelInterface_t) :: exc
  character (len=80) :: msg
  character (len=20) :: name
  ⑤ name='World'
  ⑥ call new( obj, exc )
  ⑦ call setName( obj, name, exc )
  ⑧ call sayIt( obj, msg, exc )
  ⑨ call deleteRef( obj, exc )
  print *, msg
end program helloclient
  
```

122



Handout Material: Code Notes

- ① Use statement for the greetings.English type
- ② Use statement for the sidl.BaseInterface type
- ③ Obj is a F90 derived type we get from the using statement, note the “_t” extension that prevents it from colliding with the using statement.
- ④ Exc is used to hold exceptions thrown by methods
- ⑤ In C/C++ examples, this variable would be initialized by a the command-line variable “argv[1]”, but its trickier to do portably in F90 and too long, so I just initialize the name to “World”.
- ⑥ Obj is not yet initialized. The Babel idiom in F90 is to call new() to initialize the Babel type. In other languages its _create(). NOTE: good code would add error checking.
- ⑦ setName() puts data into the obj. It sets its state.
- ⑧ saylt() returns the entire greeting including the aforementioned name.
- ⑨ deleteRef() is a subroutine that all Babel types inherit from a parent class. All Babel objects are reference counted. When there are no more outstanding references, the object is told to clean up after itself.



Working Code: “Hello World” in F90 Using a Babel Type

```

program helloclient
  use greetings_English
  use sidl_BaseInterface
  implicit none
  type(greetings_English_t) :: obj
  type(sidl_BaseInterface_t) :: exc
  character (len=80) :: msg
  character (len=20) :: name
  name='World'
  call new( obj, exc )
  call setName( obj, name, exc )
  call saylt( obj, msg, exc )
  call deleteRef( obj, exc )
  print *, msg
end program helloclient

```

Looks like a native
F90 derived type

These subroutines
were specified in the
SIDL.

Other basic subroutines
are “built in” to all Babel
types.



Question: What language is “obj” really implemented in?

```
program hellocient
  use greetings_English
  use si dl _BaseInterface
  implicit none
  type(greetings_English_t) :: obj
  type(si dl _BaseInterface_t) :: exc
  character (len=80)          : msg
  character (len=20)          :
  name=' World'
  call new( obj, exc )
  call setName( obj, name, exc )
  call sayI t( obj, msg, exc )
  call deleteRef( obj, exc )
  print *, msg
end program hellocient
```

Answer: Can't Know!

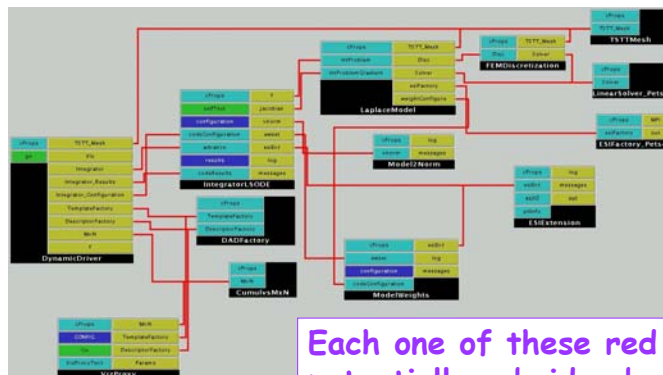
With Babel, it could be C, C++, Python, Java, Fortran77, or Fortran90/95

In fact, it could change on different runs without recompiling this code!

125



CCA uses Babel for high-performance n-way language interoperability



Each one of these red lines, is potentially a bridge between two languages. No telling which language your component will be connected to when you write it.

126



The SIDL File that defines the “greetings.English” type

```
①② package greetings version 1.0 {  
③④   interface Hello {  
⑤       void setName( in string name );  
       string sayIt ( );  
   }  
⑥   class English implements-all Hello { }  
}
```

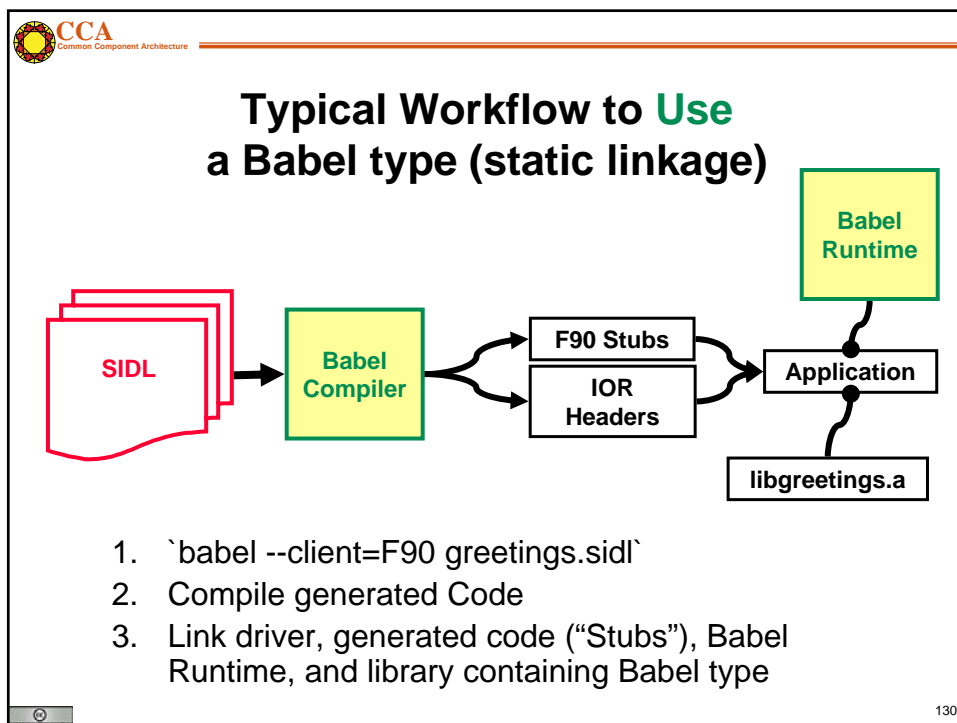
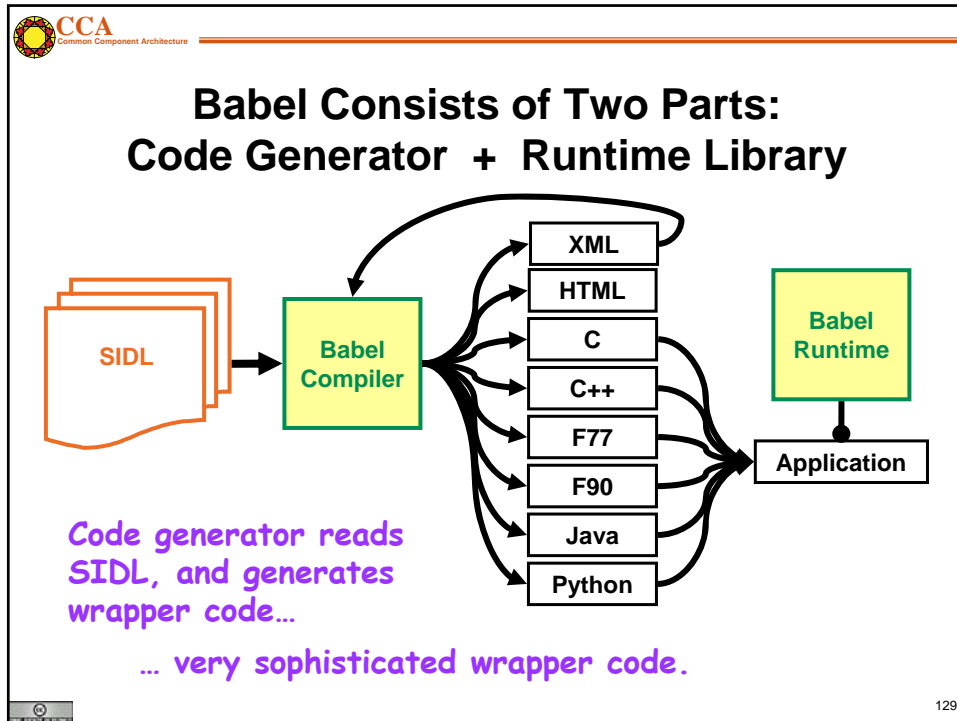
127

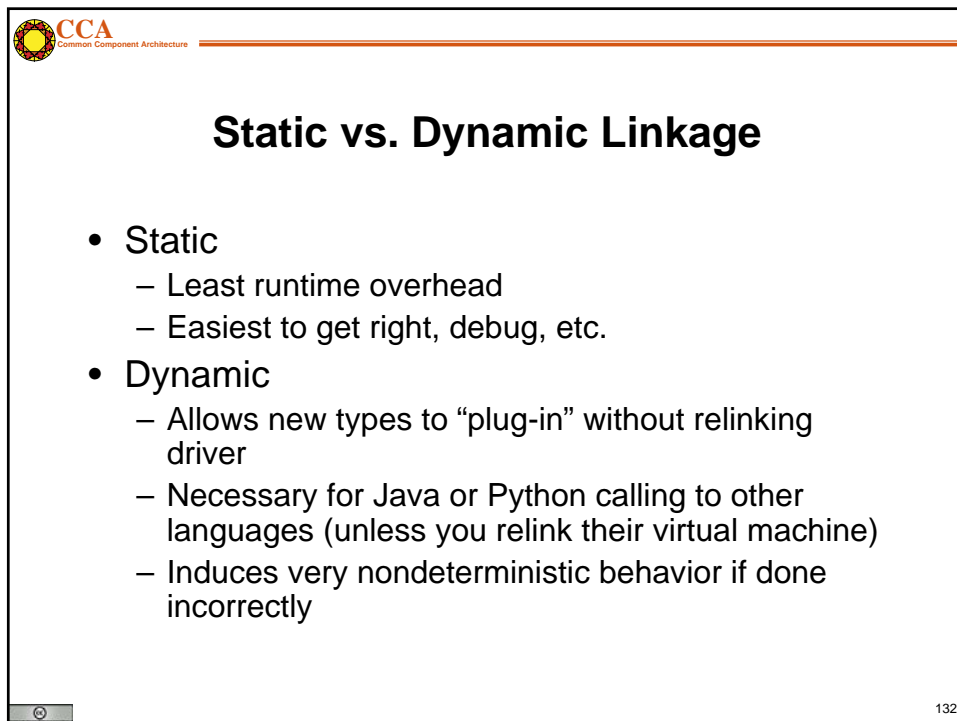
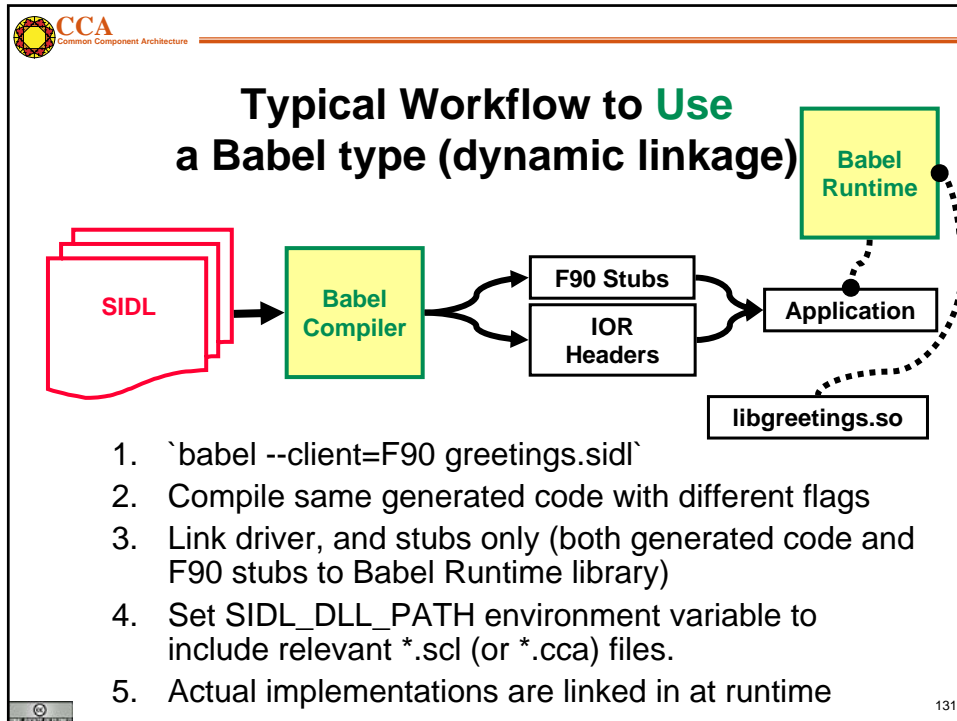


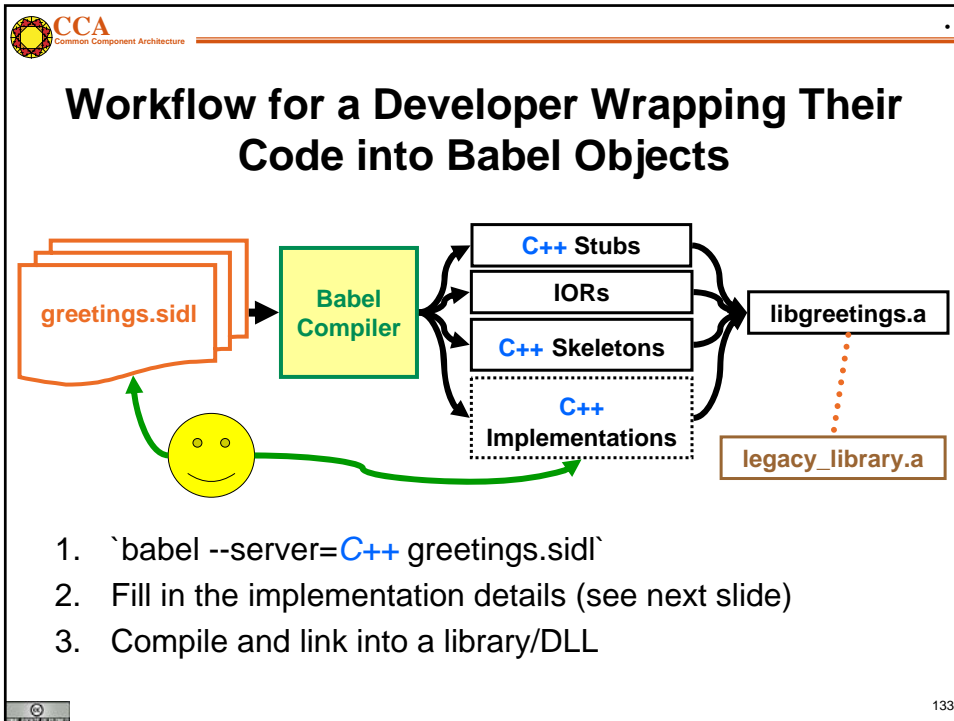
Handout Material: Code Notes

- ① Packages contain user-defined types and are used to reduce naming collisions. Packages can be nested.
- ② Packages can be versioned. User defined types must be nested inside a versioned package and gain the same version number as the innermost versioned package
- ③ SIDL has a inheritance model similar to Java and Objective C. Classes can inherit multiple interfaces, but at most one implementation (other class).
- ④ An interface describes an API, but doesn't name the implementation.
- ⑤ Note that arguments have mode, type, and name. Mode can be one of “in”, “out”, and “inout”. These SIDL modes have slightly different semantics than Fortran90 “intents”.
- ⑥ This class generates English greetings. One could imagine a strategy for internationalization that uses the Hello interface everywhere, but loads in English, French, or whatever classes based on user's preference.

128







CCA
Common Component Architecture

Implementation Details Must be Filled in Between Splicer Blocks

```
namespace greetings {
class EngIshImpl {
private:
    // DO-NOT-DELETE spl i cer. begi n(greeti ngs. EngI sh. _i mpl )
    string d_name;
    // DO-NOT-DELETE spl i cer. end(greeti ngs. EngI sh. _i mpl )

string
greeti ngs::EngI sh_i mpl::sayI t()
throw ()
{
    // DO-NOT-DELETE spl i cer. begi n(greeti ngs. EngI sh. sayI t)
    string msg("Hello ");
    return msg + d_name + "!";
    // DO-NOT-DELETE spl i cer. end(greeti ngs. EngI sh. sayI t)
}
```

134



Quick Review of Babel in general before proceeding to CCA specifics

- Babel can be used as a standalone tool
- Each language binding strikes a balance
 - support the SIDL type system (OO, exceptions, etc.)
 - provide it in a manner “natural” to experienced programmers in the target language
- For more details about Babel and SIDL
 - SC|04 tutorial slides for Babel
<http://www.llnl.gov/CASC/components/docs/sc04.html>
 - Babel User's Guide (aka. the BUG)
http://www.llnl.gov/CASC/components/docs/users_guide/



135



CCA uses SIDL to specify APIs and Type Hierarchy for Frameworks, Services, Components, & Ports

- **A CCA framework must**
 - implement gov. cca. AbstractFramework,
 - provide a gov. cca. ports. BuilderService,
 - etc.
- **A CCA port must**
 - be a SIDL interface extending gov. cca. Port
- **A CCA component must**
 - be a SIDL class implementing gov. cca. Component

The CCA Specification is a SIDL file.



136



How to write a Babelized CCA Component (1/3)

1. Define “Ports” in SIDL

- CCA Port =
 - a SIDL Interface
 - extends gov.cca.Port

```
package functions version 1.0 {  
    interface Function extends gov.cca.Port {  
        double evaluate( in double x );  
    }  
}
```



137



How to write a Babelized CCA Component (2/3)

2. Define “Components” that implement those Ports

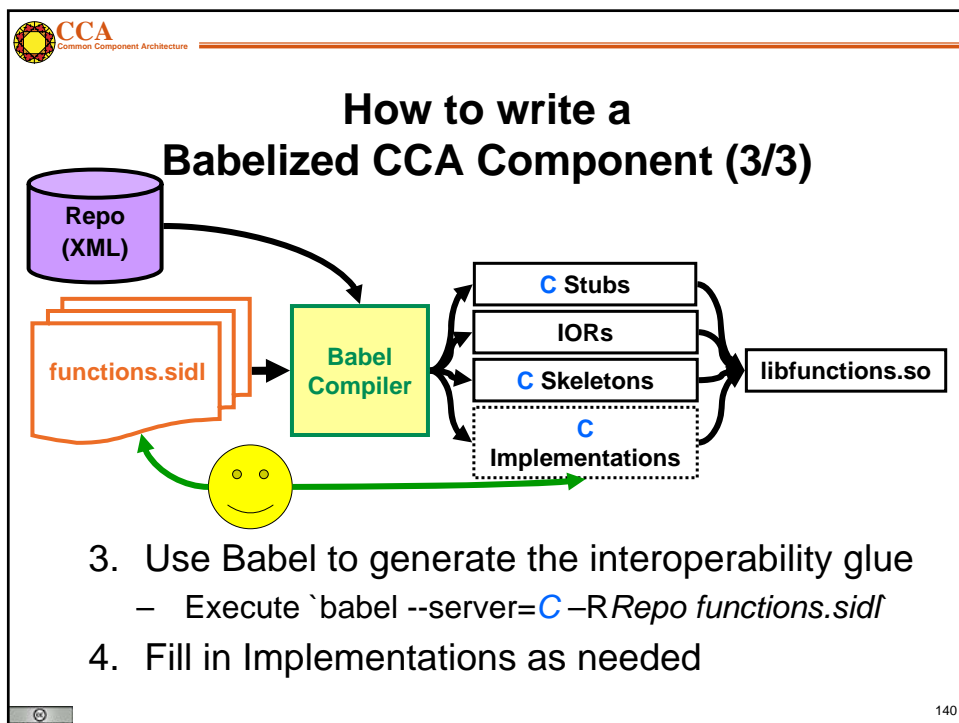
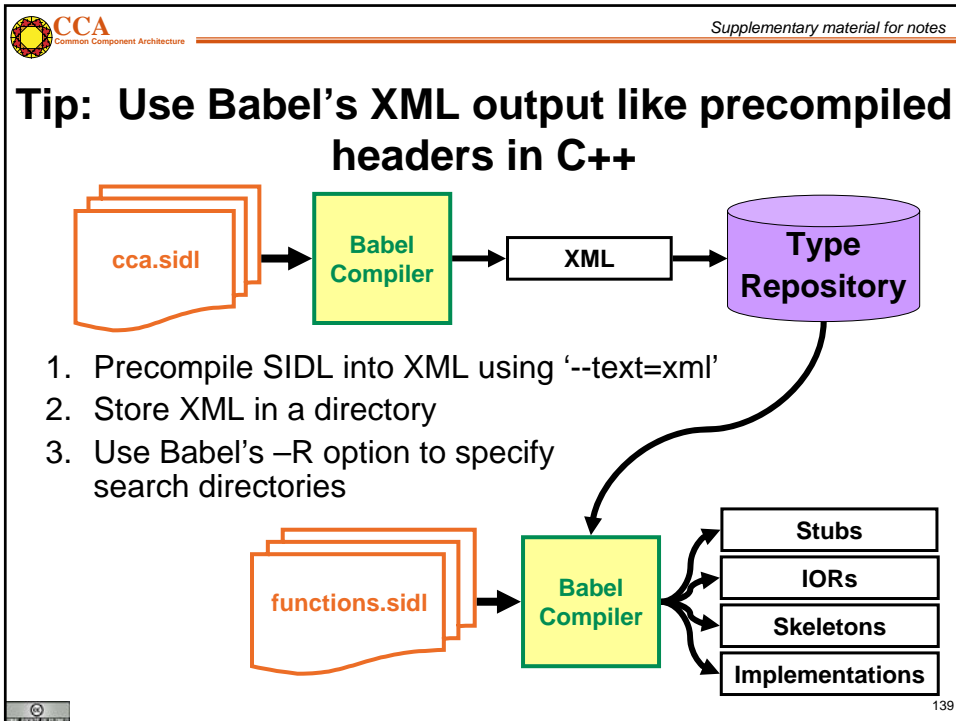
- CCA Component =
 - SIDL Class
 - implements gov.cca.Component (and any provided ports)

```
class LinearFunction implements functions.Function,  
                                gov.cca.Component {  
    double evaluate( in double x );  
    void setServices( in cca.Services svcs );  
}
```

```
class LinearFunction implements-all  
    functions.Function, gov.cca.Component { }
```



138





Review: Goal of This Module

Learn how existing code is

- Wrapped into Babel objects, and
- Promoted to CCA components

In the process, will also covered

- Babel Basics
- Scientific Interface Definition Language (SIDL)
- Code sample (“Hello World!”)



Contact Information

- Project: <http://www.llnl.gov/CASC/components>
- Project Team Email: components@llnl.gov
- Mailing Lists: majordomo@lists.llnl.gov
subscribe [babel-users](#) [email address]
subscribe [babel-announce](#) [email address]
- Bug Tracking: <https://www.cca-forum.org/bugs/babel/>
or email to babel-bugs@cca-forum.org





CCA

Common Component Architecture

Using CCA: Approaches & Experience

CCA Forum Tutorial Working Group

<http://www.cca-forum.org/tutorials/>
tutorial-wg@cca-forum.org



This work is licensed under a [Creative Commons Attribution 2.5 License](https://creativecommons.org/licenses/by/2.5/)

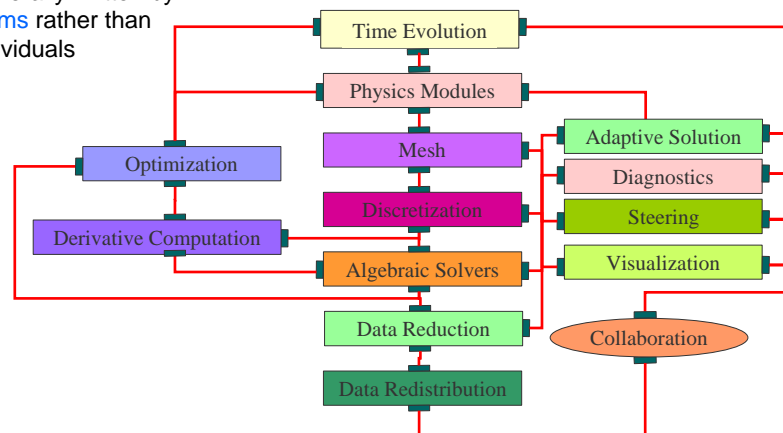
143



CCA
Common Component Architecture

Modern Scientific Software Development

- Complex codes, often **coupling** multiple types of physics, time or length scales, involving a broad range of computational and numerical techniques
- Different parts of the code require significantly **different expertise** to write (well)
- Generally written by **teams** rather than individuals



144



Using CCA to Help Manage Complexity

- Application areas participating in the CCA:
astronomy, astrophysics, biological and medical simulations, chemically reacting flow, climate and weather modelling, combustion, computational chemistry, data management, fusion and plasma physics modelling, linear algebra, materials science, molecular electronics, nanoscience, nuclear power plant simulations, structured adaptive meshes, unstructured meshes, and visualization
- Research agencies sponsoring software development using the CCA:
DOE (SciDAC, Office of Science, NNSA/ASC), NASA, NIH, NSF, DoD, European Union



Outline

- Developing Components
 - Strategies for both developing new codes and wrapping existing codes
- Case Studies
 - Chemistry project
 - Moderate-sized multi-disciplinary, multi-institutional team
 - Using Ccaffeine framework, SIDL components
 - Combustion toolkit
 - Small team, both new and wrapped codes
 - Using Ccaffeine framework, C++ components





Developing Components (Both New Codes and Wrappers to Existing Codes)

- Productivity Benefits
- Application Decomposition Strategies
- Interface Design Issues
 - Social factors
 - Technical factors
- Implementation Issues and Patterns



147



CCA Productivity Benefits

- Fast algorithmic **experiments** and benchmarks by substituting components
- Once ports are defined, domain-expert component implementers can **work separately** in their own **favorite languages**
- Work of **transient contributors** remains as well-defined, lasting components
- Wrapped legacy portions **need not be reimplemented or reverified**



148



Components in the Small: Impacts within a Project

Benefits include:

- Rapid testing, debugging, and benchmarking
- Support for implementation-hiding discipline
- Coordination of independent workers
- Interface change effects across components are clear and usually automatically found by compilers if overlooked
- Object-orientation made simpler for C and Fortran



Components in the Large: Connecting Multiple Projects

Benefits include:

- SIDL can be used to facilitate the interface consensus processes
- Different sub-projects do not have to agree on one implementation language
- Developers who never meet in person have an excellent chance of code integration working on the first try

Costs include:

- Consensus can be expensive to obtain
- Writing code for others to use is more difficult than writing it just for yourself





Application Decomposition Strategies

- Conceptually decompose the application into
 - cutting-edge areas (*less stable*)
and
 - areas that can employ existing component-based libraries (*more stable*)
- Decompose each area into components for
 - physics
 - mathematics
 - data managementas dictated by the application; sketch a typical component layout
- Many components will encapsulate *algorithmic logic only*, with little or no private data
- Most HPC applications will have a *central data abstraction* that provides data memory management and parallel communication
- In a multilanguage application, all *I/O may need to be isolated* into components written in a single common language (file based I/O should not be affected)
- Component boundaries (and port interfaces) may be set to *isolate proprietary code or difficult contributors*



151



Interface Design: Social Factors (Defining Ports to Wrap Existing Code)

- Will the port hide just one implementation, or will there need to be plug compatibility with other implementations? From other teams?
- Who defines the interface and maintains it?
 1. Project dictator? (*Fast...*)
 2. The owner of the legacy functionality? (*Slow, if not you...*)
 3. A standards committee? (*Really slow...*)
- How many iterations of redefining the ports will the customers tolerate?



152



Interface Design: Technical Factors

- Do we make a single large port look like the underlying library or divide functions into groups on separate ports?
- Should a function with many optional arguments be split into several alternative functions with simpler usage?
- Do we make the ports more general than the existing code?
- Do we require the ports to work across languages?
Across networks?
 - If not, gains in efficiency or coding ease might be had
 - If so, memory management and I/O challenges may arise



153



Implementation Issues in Wrapping

- Do we split large libraries into several components?
 - Splitting is difficult to do if global variables or common blocks are widely used.
- Do we expect more than one implementation instance of a port in a single run-time?
 - If not, interface contracts may include global side effects
- Do we integrate the wrapper code in the existing code's development and build processes?
 - If not, how do we ensure build consistency and on-going wrapper support?
- Code bases with large interfaces need automated wrapping tools
 - E.g., [see Chasm info in the Tools module](#) of the tutorial



154



Benefits of Wrapping Code Using CCA

- Setting a language-neutral interface definition (SIDL) can greatly clarify design discussions
- Provides a chance to reorganize the interface and hide globals
- Allows testing of alternate versions if doing performance studies
- Allows easy “experimentation” with new algorithms
- Software discipline is enforced, not optional
- Implementation decisions (to split libraries, etc) can be easily revised over time if interfaces remain constant (possibly with the addition of new interfaces)



155



Interface Design for New Code

- Write SIDL for each connection (port) in the sketched component layout
- If two ports must always be used together, consider merging them
- Review SIDL drafts for near-duplication of ports
- Avoid creating interface contracts that require using hidden global data
- Consider exporting tuning and/or configuration parameter inputs as a port
- All the design issues from wrapping existing code apply, also
- ***Interfaces will change.***



156



Recommended Implementation Patterns

- Expect to decompose initial components further as work progresses and requirements expand
- Build systems (i.e. `make`) should be kept as simple as possible
 - Keep a subdirectory for port definitions and any implementation-independent glue code derived from the ports
 - Keep each component (and any wrapped code) in its own subdirectory
 - Keep application-wide flags in a configure script or an include file common to all components and ports
 - Consistency is key. Extract build flags from `cca-spec-babel-config` and if possible compile & link with `babel-libtool`



157



Outline

- Developing Components
 - Strategies for both developing new codes and wrapping existing codes




Case Studies

- Chemistry project
 - Moderate-sized multi-disciplinary, multi-institutional team
 - Using Ccaffeine framework, SIDL components
- Combustion toolkit
 - Small team, both new and wrapped codes
 - Using Ccaffeine framework, C++ components




158





CCA
Common Component Architecture

Case Study 1: Chemistry Project

- Funded via SciDAC initiative
- Initial focus: Full-featured components for structure optimization
 - **Chemistry models** provided by MPQC (SNL) & NWChem (PNNL)
 - **Numerical optimization** provided by TAO (ANL) solvers
 - **Linear algebra** provided by GA (PNNL) and PETSc (ANL)
- Recent work:
 - Multi-level parallelism
 - Low-level chemistry model integration (e.g., molecular integrals)




159



CCA
Common Component Architecture

CCA-Chemistry Project Participants

Pacific Northwest National Laboratory


Theresa L. Windus
Yuri Alexeev
Manojkumar Krishnan
Jarek Nieplocha
Carl Fahlstrom
Elizabeth Jurrus


Argonne National Laboratory

Steve Benson
Lois Curfman McInnes
Jason Sarich


Sandia National Laboratory

Curtis L. Janssen
Joseph P. Kenney







Pacific Northwest
National Laboratory
Operated by Battelle for the
U.S. Department of Energy



ARGONNE NATIONAL LABORATORY
UNIVERSITY OF CHICAGO



Sandia
National
Laboratories


160



CCA Impacts in Computational Chemistry

Through 4 chemistry applications we consider different facets of CCA functionality:


- **Molecular Geometry Optimization**
 - Combining diverse expertise of 5 different research groups
- **Lennard-Jones Molecular Dynamics**
 - Achieving good scalability and low CCA overhead
- **Multi-level Parallelism in Computational Chemistry**
 - Combining SPMD and MPMD parallelism
- **Molecular Integral Evaluation**
 - Component interoperability at deeper levels within chemistry libraries



161



App 1: Molecular Geometry Optimization: Combining Diverse Expertise


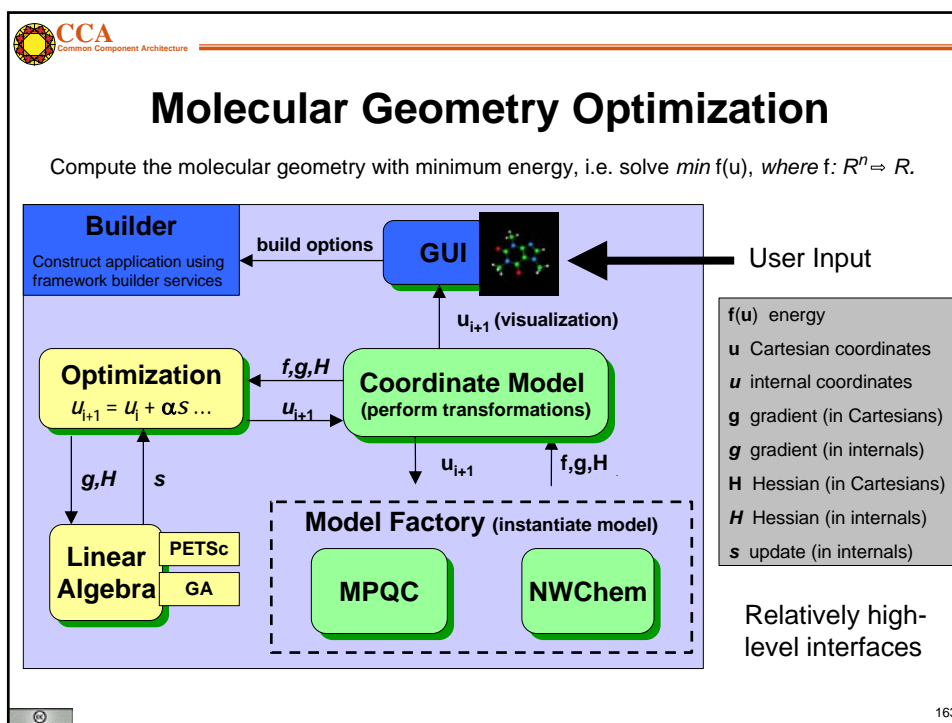


CCA quantum chemistry application using components based on:

- MPQC, NWChem (chemistry – energy evaluation)
- GA, PETSc (parallel data management and linear algebra)
- TAO (numerical optimization)



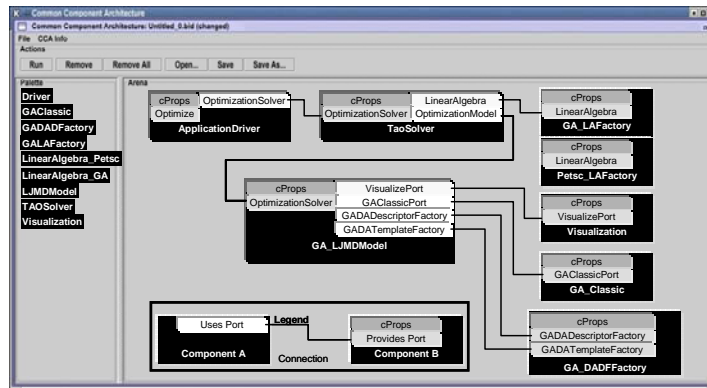
162





Common Component Architecture

App 2: Molecular Dynamics




CCA component wiring diagram for Lennard-Jones energy optimization, which integrates TAO, GA and Lennard-Jones molecular dynamics components

165




Common Component Architecture

Good CCA Scalability, Negligible Overhead



Speedup in the Lennard-Jones potential energy optimization for 32,768 and 65,536 particles

Overhead of CCA component (using SIDL) vs. "native" C++ implementation of a parallel Lennard-Jones molecular dynamics simulation



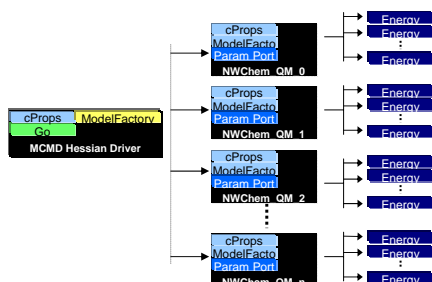
% Overhead using CCA Components

166



App 3: Multilevel Parallelism Using CCA

- Massive numbers of CPUs in future systems require leading edge tools to exploit all available parallelism
- GA helps exploit multi-level parallelism
- Multi-level parallelism using CCA & GA processor groups: Combining SCMD (Single Component Multiple Data) and MCMD (Multiple Component Multiple Data) paradigms




The MCMD Driver launches multiple instances of NWChem components on subsets of processors (also assigns a GA communicator for every instance). Each NWChem QM component does multiple energy computations on subgroups.

167



Multilevel Parallelism - MCMD

NWChem Scalability of Numerical Hessian Computation on 11.8 TFlop/s HP Cluster at PNNL



M. Krishnan, Y. Alexeev, T. Windus, and J. Nieplocha, Multilevel Parallelism in Computational Chemistry using the Common Component Architecture and Global Arrays, Proceedings of the IEEE/ACM SC2005 Conference.

168

CCA
Common Component Architecture

App 4: Molecular Integrals: Deeper Levels of Interoperability among Chemistry Packages

- As a first step toward low-level integration of chemistry models, capabilities to provide and utilize molecular integral evaluator components are being added to MPQC and NWChem (also working with GAMESS developers)
- Specialized integral capabilities are provided to CCA chemistry models
 - Relativistic integrals in NWChem
 - Non-standard 2-electron integrals for linear-R12 theories in MPQC

169

CCA
Common Component Architecture

CCA Impacts in Computational Chemistry: Review

Through 4 chemistry applications we considered different facets of CCA functionality:


- Combining diverse expertise of 5 different research groups
- Achieving good scalability and low CCA overhead
- Implementing multi-level parallelism by combining SPMD and MPMD paradigms
- Addressing component interoperability at deeper levels within chemistry libraries

170



Case Study 2: Combustion Project

- Computational Facility for Reacting Flow Science (CFRFS)
 - <http://cfrfs.ca.sandia.gov>
 - Funded via SciDAC initiative (PI: H. Najm)
- Focus: A toolkit to perform simulations of lab-sized unsteady flames
 - Solve the Navier-Stokes w/detailed chemistry
 - Various mechanisms up to ~50 species, 300 reactions
- Consequently:
 - Disparity of **length scales** :
 - use structured adaptively refined meshes
 - Disparity of **time scales** (transport versus chemistry) :
 - use an operator-split construction and solve chemistry implicitly
 - adaptive chemistry: use **computational singular perturbation** to find and follow low dimensional chemical manifolds



J. Ray, S. Lefantzi, J. Lee, C. Kennedy, W. Ashurst, K. Smith, M. Liu, N. Trebon

171



Why Use CCA in the CFRFS Toolkit?

- Separate clearly the physics models, numerical algorithms, and the “CS” parts of the toolkit
 - Strictly functional!
- Realize the separation in software
- Tame **software** complexity
- Separate contributions by transient contributors
 - Form the bulk of the developers
- Create “chunks” of well-defined functionality that can be developed by experts (usually numerical analysts and combustion researchers)

172




Design Principles of the Toolkit / 1

- **Principal Aim: Reduce software complexity**
 - We can deal with the rest
- Functional decomposition into components
 - “Data Object” and Mesh components
 - (Large) set of numerical algorithmic components (integrators, linear/nonlinear solvers, etc.)
 - (Large) set of physical models components (gas-phase combustion chemistry, thermodynamics, fluid dynamic quantities, e.g. viscous stress tensor)
 - Handful of adaptors



Design Principles of the Toolkit / 2

- Decomposition reflected in the port design and implementation
 - Most re-implemented port is the one that exchanges a rectangular sub-domain's data for processing by components
- Sparse connectivity between components
 - i.e., components communicate with a few others
 - Large apps (component assemblies) are composed by assembling smaller, largely independent sub-assemblies
 - Sub-assemblies usually deal with a certain physics
 - Intuitive way to assemble a *multiphysics* code



CCA
Common Component Architecture

The Code

Transport subassembly

Diffusion coefficients

Chemistry reaction subassembly

Separate component subsystems for transport (black) and for reaction (orange) in a reaction-diffusion code. They two are coupled at a relatively high level.

175

CCA
Common Component Architecture

Has the Toolkit Approach Helped Tame Software Complexity?

- How has the code evolved?
 - How often have new ports been added?
 - How many rewrites have been done?
- How large are the components?
- How many ports do they have?
 - How large are the ports?
- How many ports exist?
 - i.e., Is the design general enough to support many implementations?
- What is the connectivity of components in application codes?


176



Common Component Architecture

CFRFS Toolkit Status

- Started in 2001
- 61 components today, all peers, independent, mixed and matched for combustion and shock hydrodynamics
- 7 external libraries
- Contributors: 9 in all, including 3 summer students
- Only 2 of the 9 contributors are at SNL today




A Fitzhugh-Nagumo equation being solved on a block-structured adaptively refined mesh. The top image illustrates Runge phenomena at coarse-fine interfaces (dashed ovals) when using high-order schemes (6th order interpolations with 4th order discretizations). Filtering them with an 8th order filter removes them (bottom).

177





Common Component Architecture

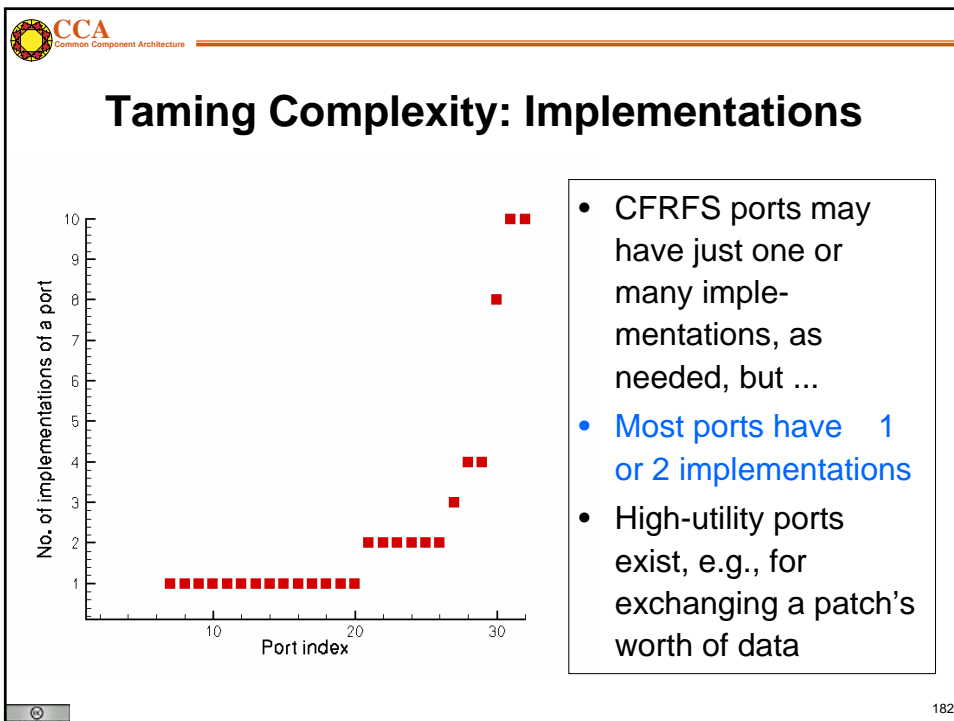
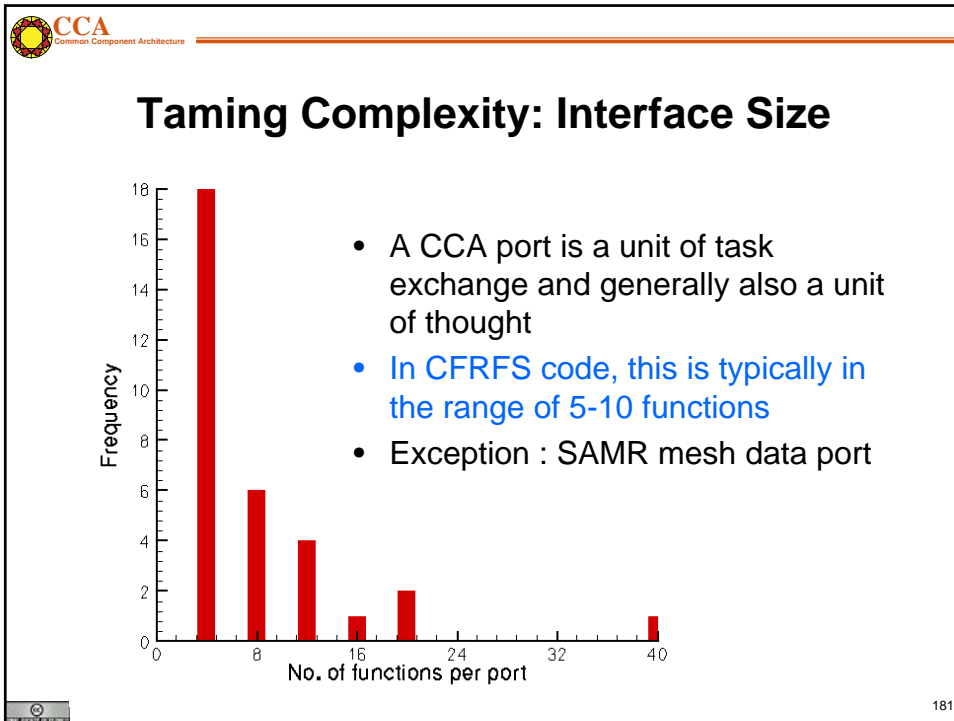
Scalability: Capability Growth without Rewrites




- Port designs typically occur in spurts followed by long component development times.
- Ports may have multiple implementations; hence the number of ports is typically less than the number of components.
- As the toolkit has matured, the number of ports is seen to be asymptoting to a slow growth rate.

178









Adaptive Codes

- CCA code is an assembly of components
 - Components are shared libraries
 - Component applications can be non-optimal because
 - The mathematical characteristics of the simulation are different from those of the component
 - The component is badly implemented
- So, can one dynamically “re-arrange” an assembly to improve performance?
- Simplistically, 2 approaches:
 - Create performance model per component, use best one
 - Create expert system that analyzes problem and picks good solution strategy



185



Performance Measurement in a Component World

- Different from traditional software
 - The researcher may not be the author of the component
 - Will need performance info at component granularity
 - Will need a non-intrusive way of getting it
- Proxies: Simple components to be inserted between a caller and a callee component
 - Intercepts and forwards method calls
 - Use to log size of arrays, calling parameters, counters, etc.
 - Can be used to turn on a “clock”; turned off when returning to caller
- Proxy’s ports are the same as the callee’s
 - Functionality is pretty mundane ~ mostly pass through...
 - Can it be automatically generated?



186



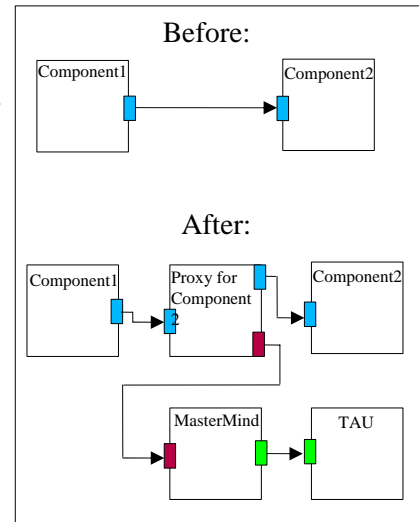
Integrated Performance Measurement

Measurement infrastructure:

- Proxy
 - Generated automatically using PDT
 - Logs parameters to MasterMind
- MasterMind
 - Collects and stores all measurement data
- TAU
 - Makes all performance measurements



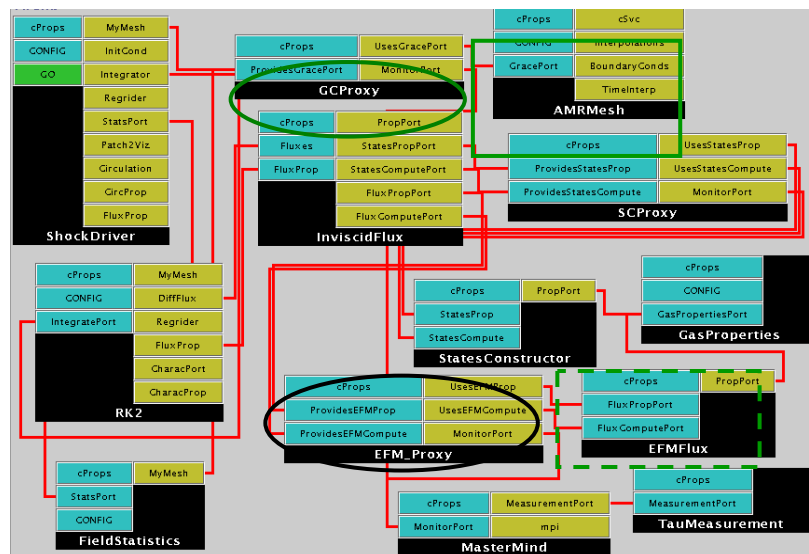
Work done at University of Oregon
by A. Malony and team



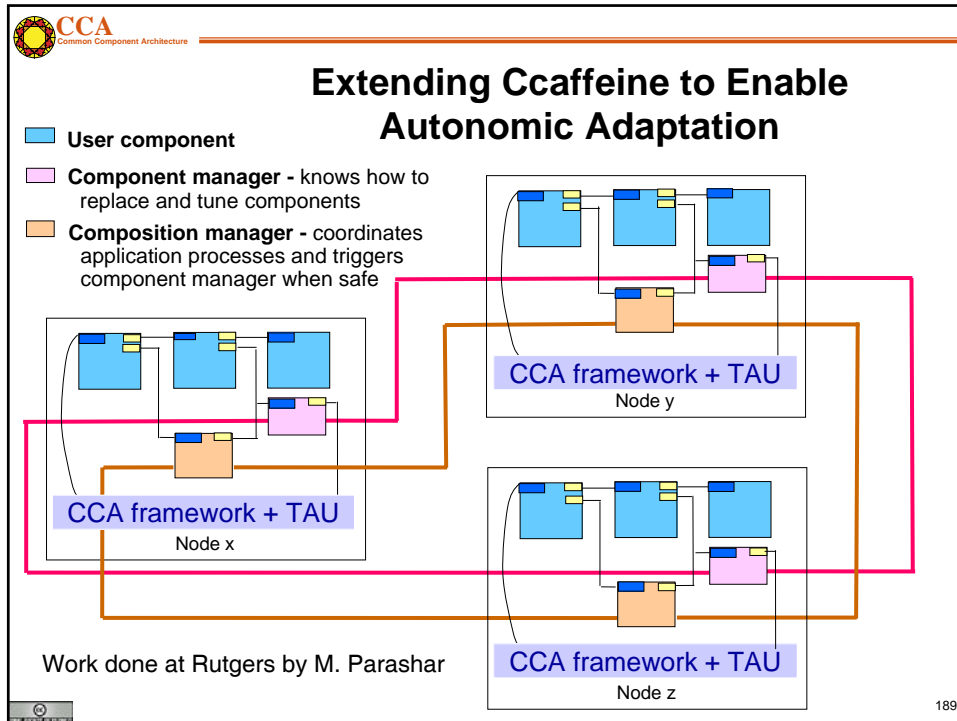
187



Component Application with Proxies



188



CCA
Common Component Architecture

Incorporating Other Frameworks

- Chombo (by P. Colella, LBNL) has solvers (AMRGodunov, AMRElliptic, etc.) that:
 - Work on block structured adaptive meshes
 - Accept Chombo-specific data structures
 - But fundamentally require:
 - A double pointer, an array, where variables on a patch (box) are stored in blocked format
 - Bunch of integer arrays that describe the array
- Challenge: Can Chombo be used within CCA?
- Need a standardized way of getting data into Chombo
 - Pointer aliasing, not data copy
- Implementation of this “standard” interface

Using Chombo to solve the Poisson equation (needed for CFRFS flame simulations).

190



CCA
Common Component Architecture

Using CCA: Summary

- Review of guidelines for developing high-performance scientific components (both new code and wrappers for existing code)
- CCA is an enabling technology for scientific applications
 - Has enabled mathematicians, chemists, combustion scientists, and computer scientists to contribute new strategies that are shrink-wrapped for easy re-use
 - Apart from **science research**, also spawned **new research directions in CS**
 - Has enabled research scientists to investigate unconventional approaches, for example multilevel parallelism and dynamic adaptivity
- For more info on the CCA applications/case studies, see:
 - Chemistry: <http://www.cca-forum.org/~cca-chem>
 - Combustion: <http://cfrfs.ca.sandia.gov>
- **Different facets of CCA components may be useful within different projects ... What are **your** needs and priorities?**



191



CCA
Common Component Architecture

A Few Notes in Closing

CCA Forum Tutorial Working Group
[http://www.cca-forum.org/tutorials/
tutorial-wg@cca-forum.org](http://www.cca-forum.org/tutorials/tutorial-wg@cca-forum.org)



This work is licensed under a [Creative Commons Attribution 2.5 License](http://creativecommons.org/licenses/by/2.5/)

192



Resources: Its All Online

- Information about all CCA tutorials, past, present, and future:
<http://www.cca-forum.org/tutorials/>
- Specifically...
 - Latest versions of hands-on materials and code:
<http://www.cca-forum.org/tutorials/#sources>
 - Hands-On designed for self-study as well as use in an organized tutorial
 - Should work on most Linux distributions, less tested on other unixen
 - Still evolving, so please contact us if you have questions or problems
 - Archives of all tutorial presentations:
<http://www.cca-forum.org/tutorials/archives/>
- Questions...
help@cca-forum.org



Getting Help

- We want to help insure you have a good experience with CCA, so let us know if you're having problems!
- Tutorial or "start-up" questions
 - help@cca-forum.org
- Problems with specific tools
 - *check documentation for updated contact info*
 - cca-tools bundle (includes Chasm, Babel, Ccaffeine): [Rob Armstrong, rob@sandia.gov](mailto:Rob.Armstrong@sandia.gov)
 - Chasm: [Craig Rasmussen, crasmussen@lanl.gov](mailto:crasmussen@lanl.gov)
 - Babel: babel-users@lanl.gov
 - Ccaffeine: ccafe-users@cca-forum.org
- General questions, or not sure who to ask?
 - help@cca-forum.org





Common Component Architecture

CCA is Interactive

- Collectively, CCA developers and users span a broad range of scientific interests.
 - There's a good chance we can put you in touch with others with relevant experience with CCA
- CCA Forum Quarterly Meetings
 - Meet many CCA developers and users
 - <http://www.cca-forum.org/meetings/>
- “Coding Camps”
 - Bring together CCA users & developers for a concentrated session of coding
 - Held as needed, typically 3-5 days
 - May focus on a particular theme, but generally open to all interested participants
 - If you're interested in having one, *speak up* (to individuals or cca-forum@cca-forum.org)
- Visits, Internships, etc.



195



Common Component Architecture

Acknowledgements: Tutorial Working Group

- **People:** Benjamin A. Allan, Rob Armstrong, David E. Bernholdt, Randy Bramley, Tamara L. Dahlgren, Lori Freitag Diachin, Wael Elwasif, Tom Epperly, Madhusudhan Govindaraju, Ragib Hasan, Dan Katz, Jim Kohl, Gary Kurfert, Lois Curfman McInnes, Alan Morris, Boyana Norris, Craig Rasmussen, Jaideep Ray, Sameer Shende, Torsten Wilde, Shujia Zhou
- **Institutions:** ANL, Binghamton U, Indiana U, JPL, LANL, LLNL, NASA/Goddard, ORNL, SNL, U Illinois, U Oregon
- **Computer facilities** provided by the Computer Science Department and University Information Technology Services of Indiana University, supported in part by NSF grants CDA-9601632 and EIA-0202048.



196



Acknowledgements: The CCA

- **ANL** –Steve Benson, Jay Larson, Ray Loy, Lois Curfman McInnes, Boyana Norris, Everest Ong, Jason Sarich...
 - **Binghamton University** - Madhu Govindaraju, Michael Lewis, ...
 - **Indiana University** - Randall Bramley, Dennis Gannon, ...
 - **JPL** – Dan Katz, ...
 - **LANL** - Craig Rasmussen, Matt Sotille, ...
 - **LLNL** – Tammy Dahlgren, Lori Freitag Diachin, Tom Epperly, Scott Kohn, Gary Kurfert, ...
 - **NASA/Goddard** – Shujia Zhou
 - **ORNL** - David Bernholdt, Wael Elwasif, Jim Kohl, Torsten Wilde, ...
 - **PNNL** - Jarek Nieplocha, Theresa Windus, ...
 - **SNL** - Rob Armstrong, Ben Allan, Lori Freitag Diachin, Curt Janssen, Jaideep Ray, ...
 - **Tech-X Corp.** – Johan Carlsson, Svetlana Shasharina, Ovsei Volberg, Nanbor Wang
 - **University of Oregon** – Allen Malony, Sameer Shende, ...
 - **University of Utah** - Steve Parker, ...
- and many more... without whom we wouldn't have much to talk about!

197



Thank You!

Thanks for attending this tutorial

We welcome feedback and questions

198